

Image processing

Jan. 28, 2025

Today's Class

- Project 2
- Image filtering
- Template matching
- Edge detection

Sources of material and examples

- Image operations and filtering is a big topic. We are only scratching the surface
- We will follow Ch.11 from Corky et al (2023) textbook.
- Please read Szeliski 3.1 and 3.2
- Matlab vision toolbox
 - Make sure that any code you write that reference any of these functions has the proper citation as a comment.
 - Corke has a MV toolbox that contains most of the functions in his book. You can download from <https://github.com/petercorke/RVC3-MATLAB>

IMAGE FILTERING

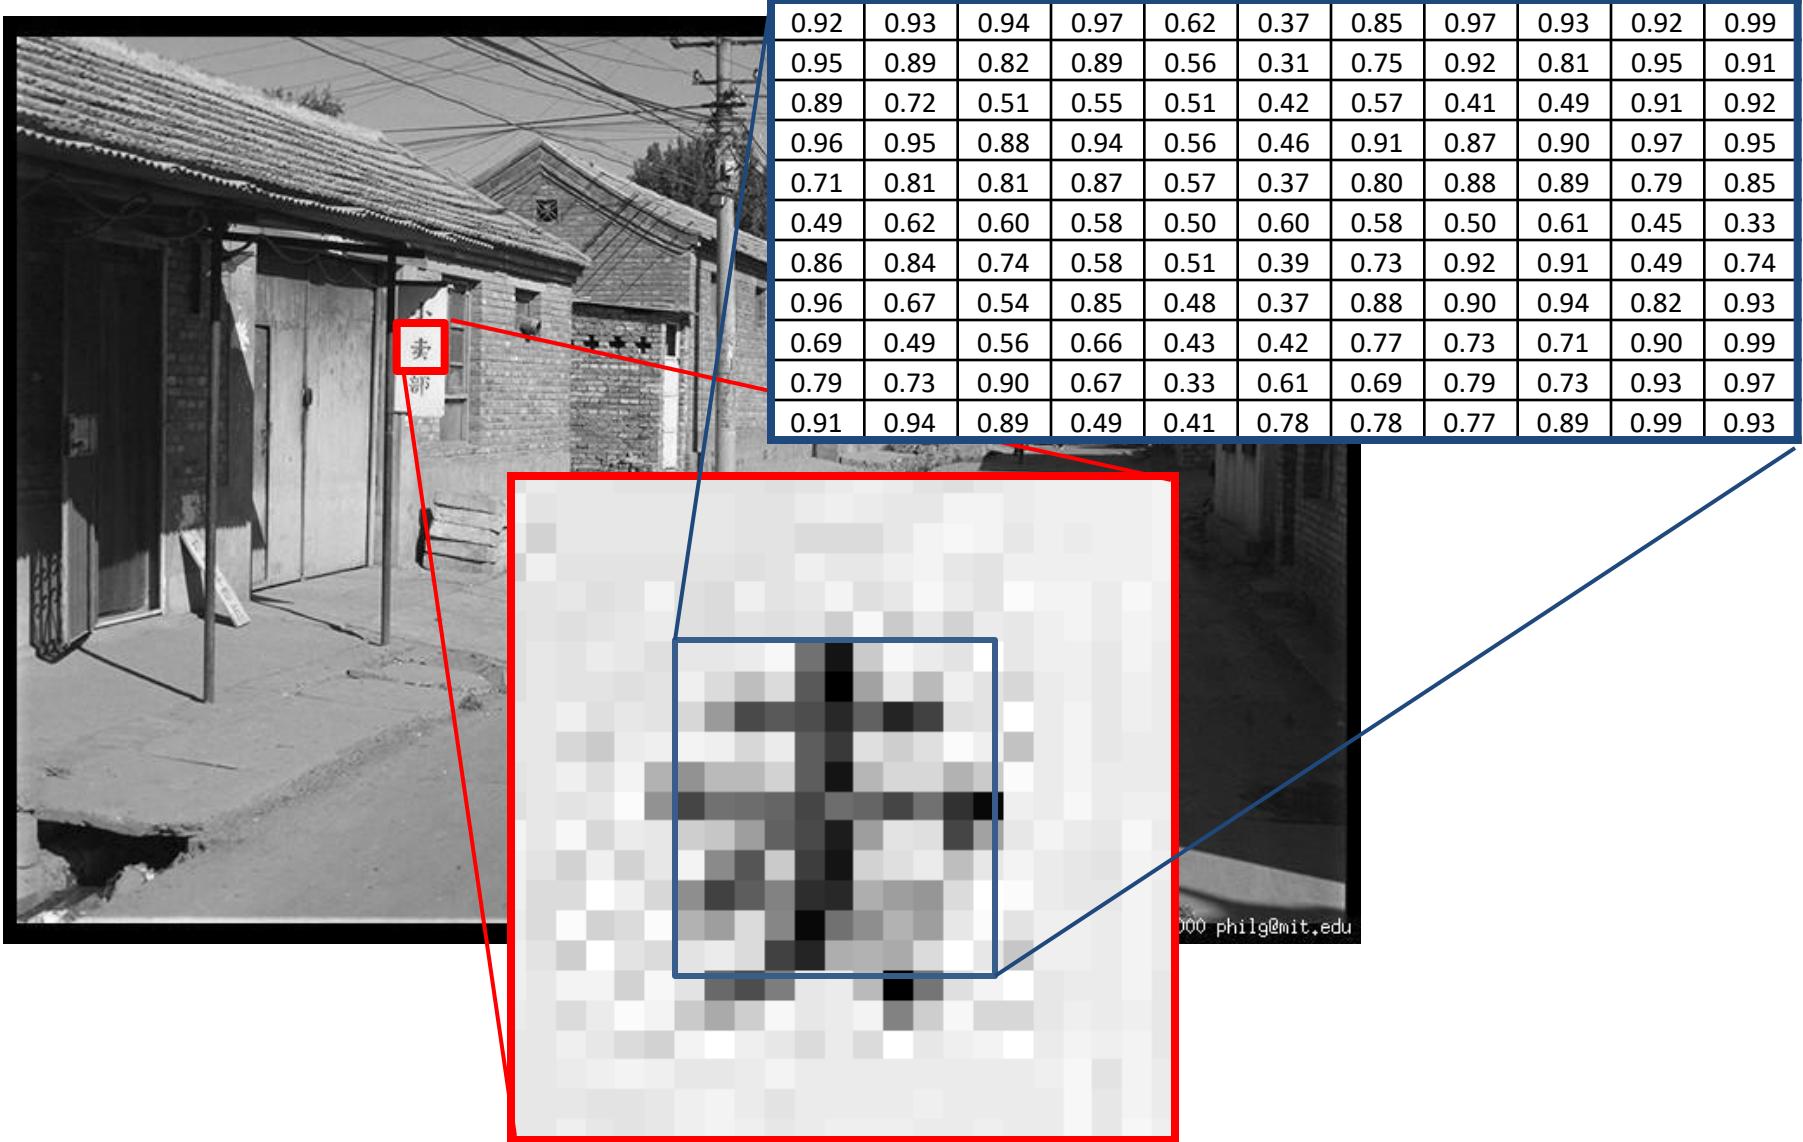
Images in Matlab

- Images represented as a matrix
- Suppose we have a NxM RGB image called “im”
 - $\text{im}(1,1,1)$ = top-left pixel value in R-channel
 - $\text{im}(y, x, b)$ = y pixels down, x pixels to right in the bth channel
 - $\text{im}(N, M, 3)$ = bottom-right pixel in B-channel
- `imread(filename)` returns a uint8 image (values 0 to 255)
 - Convert to double format (values 0 to 1) with `im2double`

The diagram illustrates a 10x10 matrix of image data. A vertical blue arrow labeled "row" points downwards, indicating the horizontal axis. A horizontal blue arrow labeled "column" points to the right, indicating the vertical axis. To the right of the matrix, three smaller 4x4 matrices are shown, labeled "R", "G", and "B", representing the Red, Green, and Blue color channels respectively. The main matrix contains numerical values ranging from 0.49 to 0.95. The "R" channel matrix shows values for the first four columns of the main matrix. The "G" channel matrix shows values for the next four columns. The "B" channel matrix shows values for the last two columns.

0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93
0.89	0.73	0.56	0.56	0.49	0.49	0.72	0.77	0.79	0.71	0.90
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93
0.89	0.73	0.56	0.56	0.49	0.49	0.72	0.77	0.79	0.71	0.90
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93

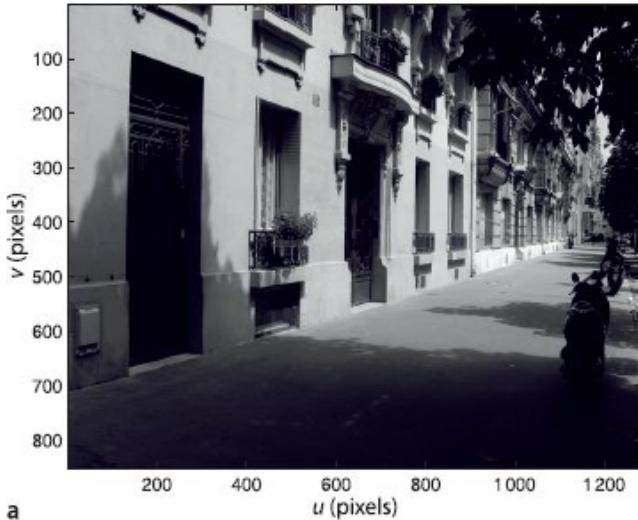
Back to grayscale intensity



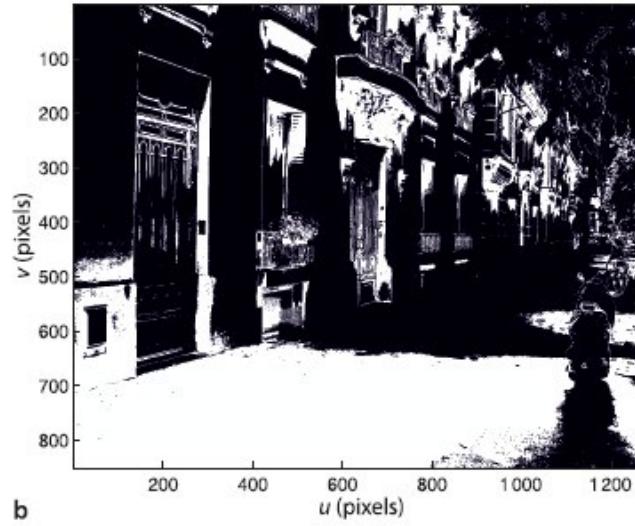
Monadic operations

- Operations on one pixel at a time
 - $O[u, v] = f(I[u, v]), \forall (u, v) \in I$
 - Examples
 - Color to grey imono(color_image_file)
- Useful in many situations
 - Histogram thresholding
 - Histogram normalization

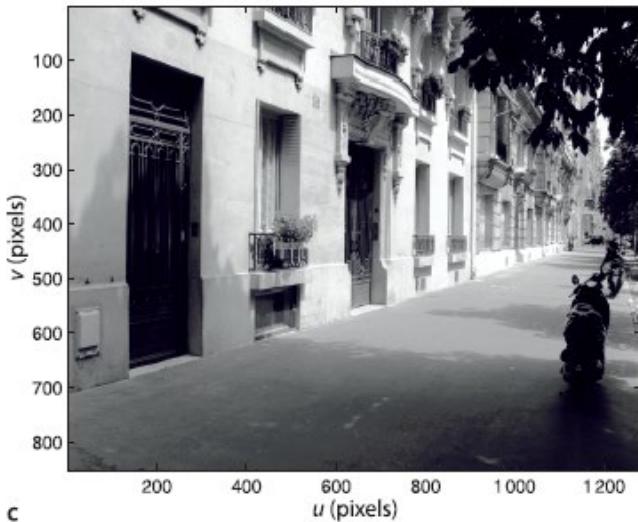
Example of monadic operations



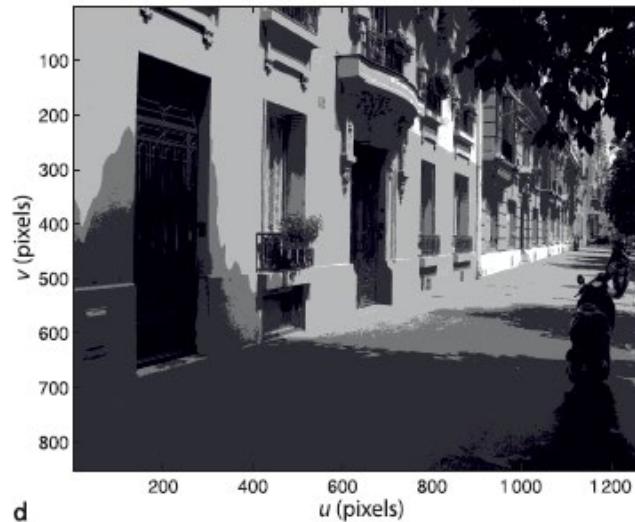
a



b



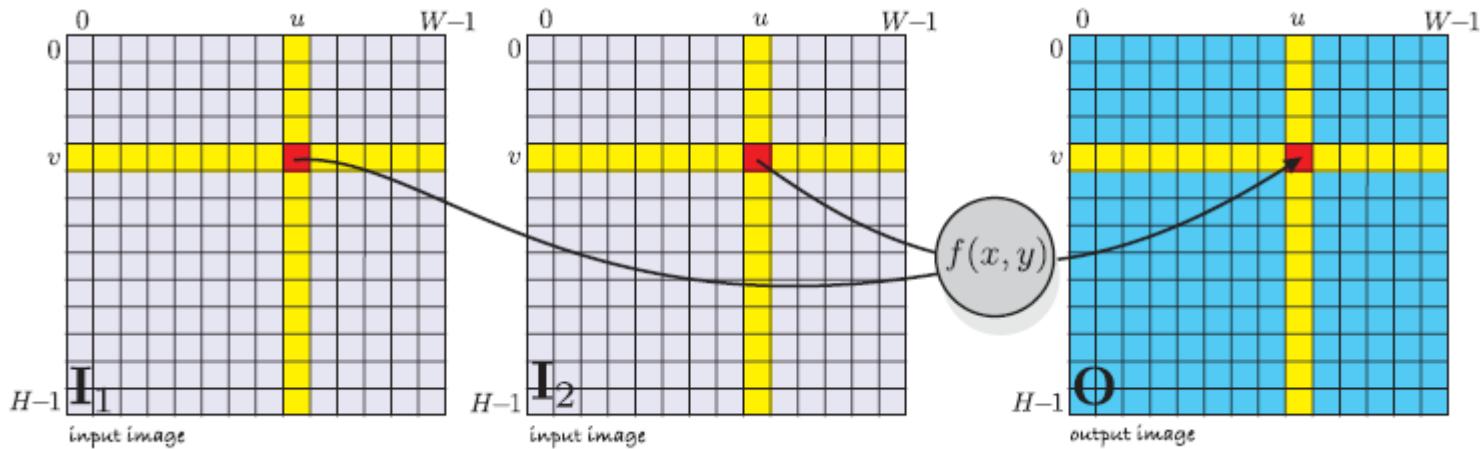
c



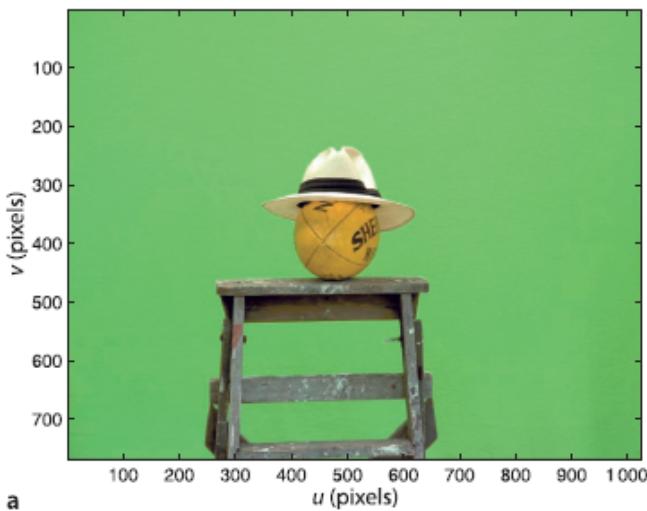
d

Diadic Operations

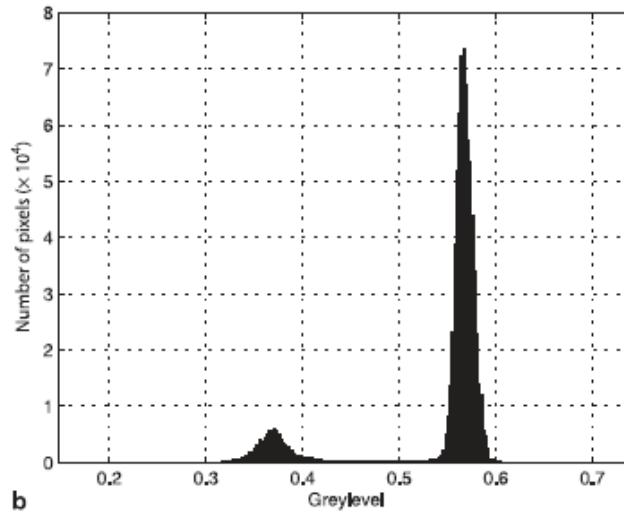
- Operations on two images to result in one image
 - $O[u, v] = f(I_1[u, v], I_2[u, v]), \forall (u, v) \in I_i$



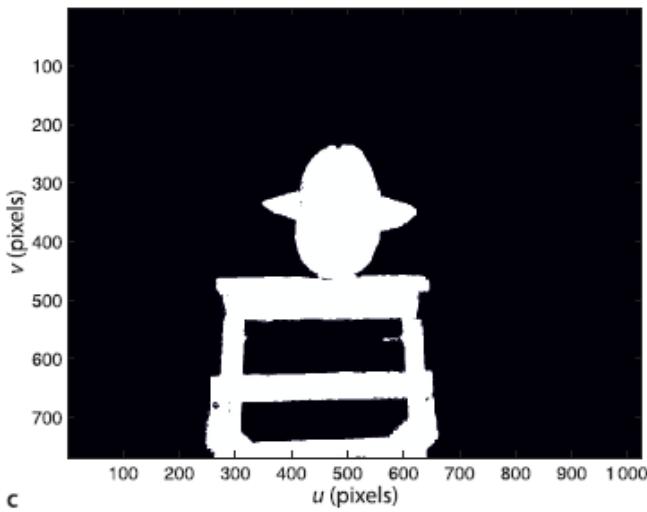
Example of diadic operation



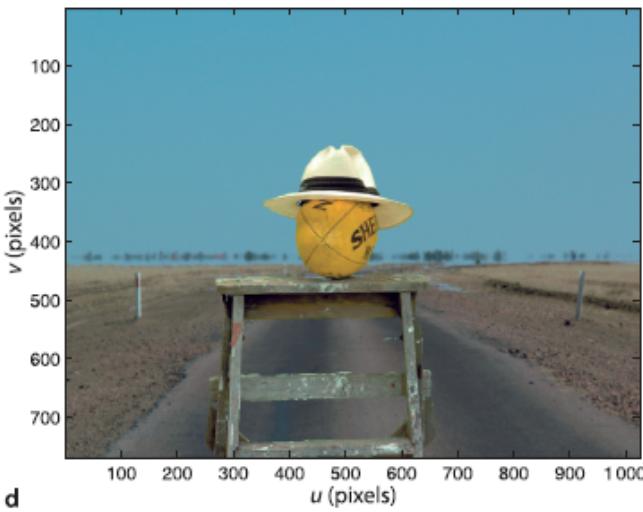
a



b



c



d

Extracting background

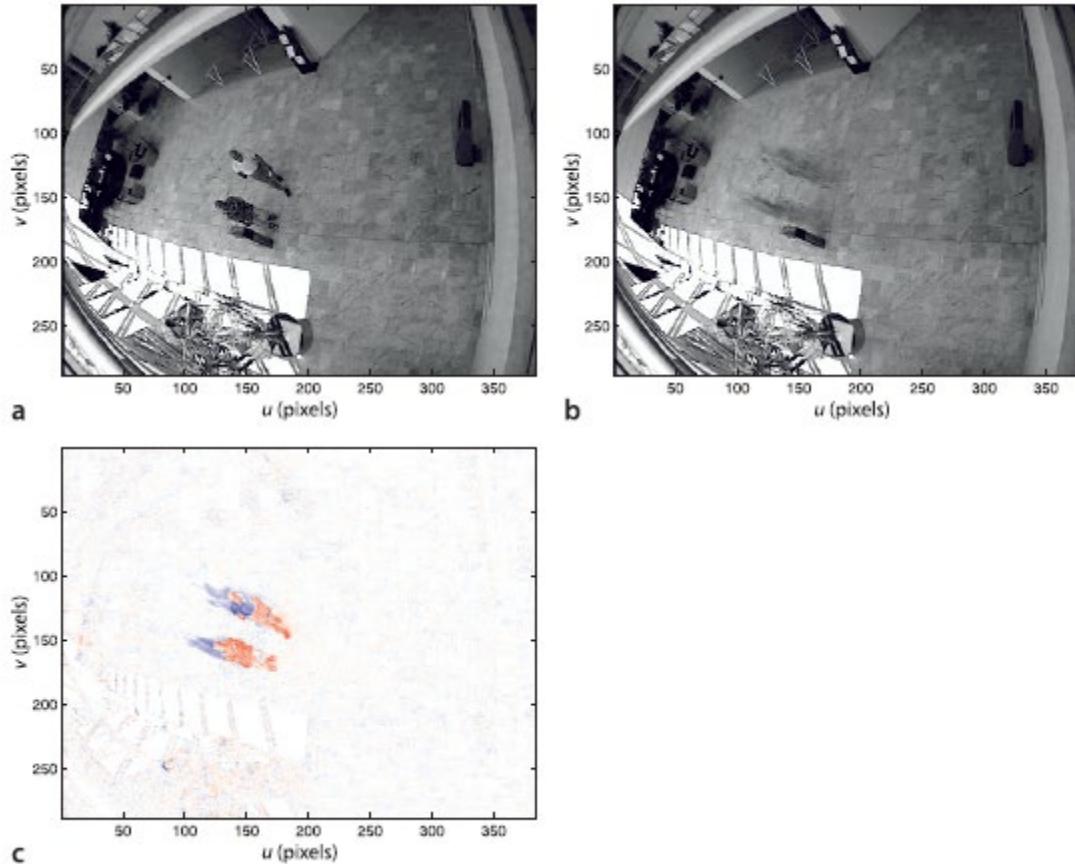


Image filtering

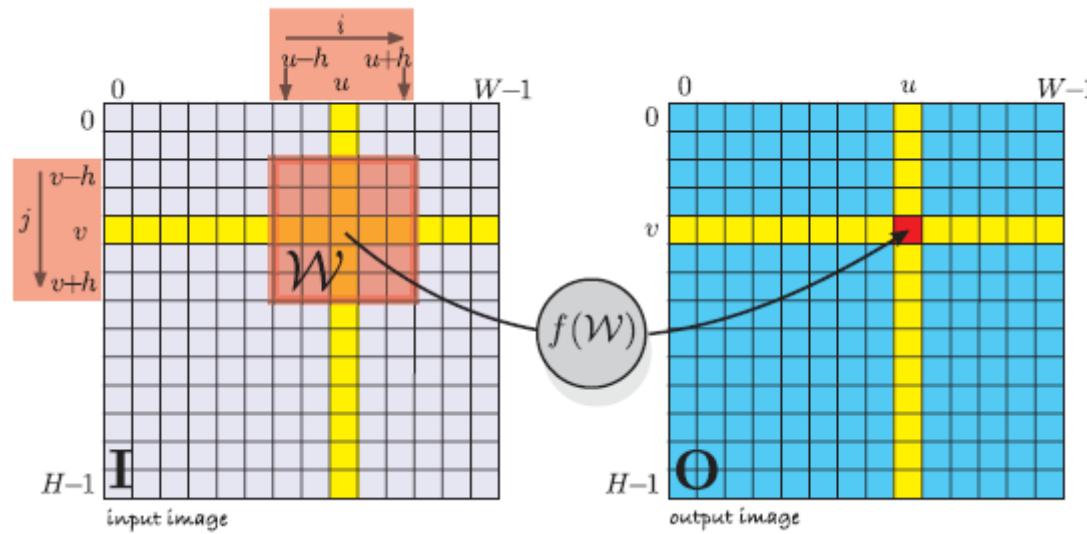
- Image filtering: compute function of local neighborhood at each position
- Really important!
 - Enhance images
 - Denoise, resize, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching

Type of filters

- Image filters in spatial domain
 - Filter is a mathematical operation of a grid of numbers
 - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
 - Filtering is a way to modify the frequencies of images
 - Denoising, sampling, image compression
- Templates and Image Pyramids
 - Filtering is a way to match a template to the image
 - Detection, coarse-to-fine registration

Spatial operations (filters)

- $O[u, v] = f(I[u + i, v + j]), \forall (i, j) \in W, \forall (u, v) \in I$
Where W is the window, $w \times w$ square region



$f(\cdot)$ can be linear and non-linear functions

Example: box filter

$g[\cdot, \cdot]$

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

Image filtering

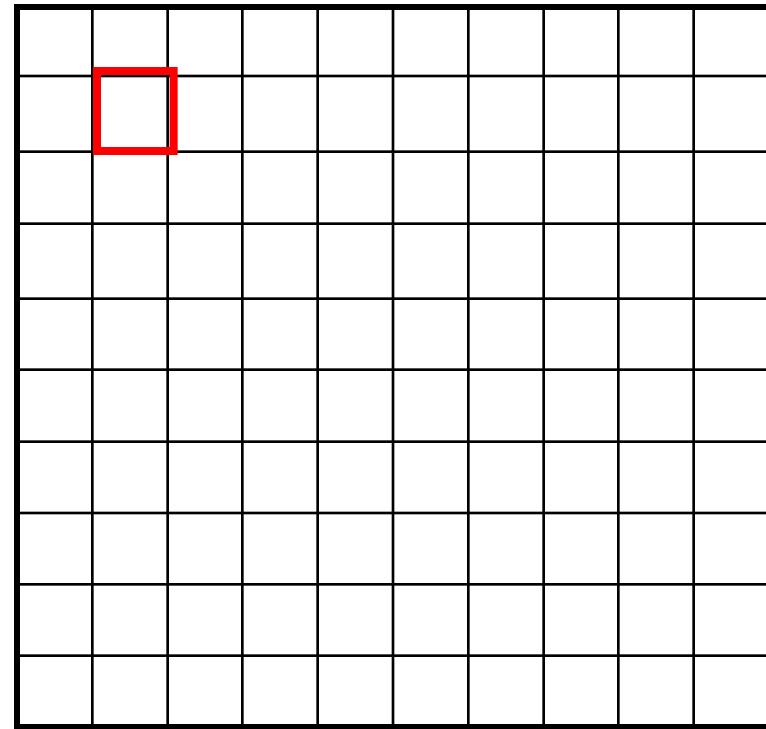
$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$



$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

Credit: S. Seitz

Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10									

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0 10 20

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

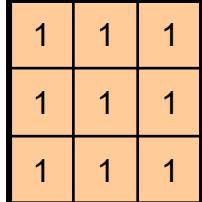
$h[.,.]$

0 10 20 30

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

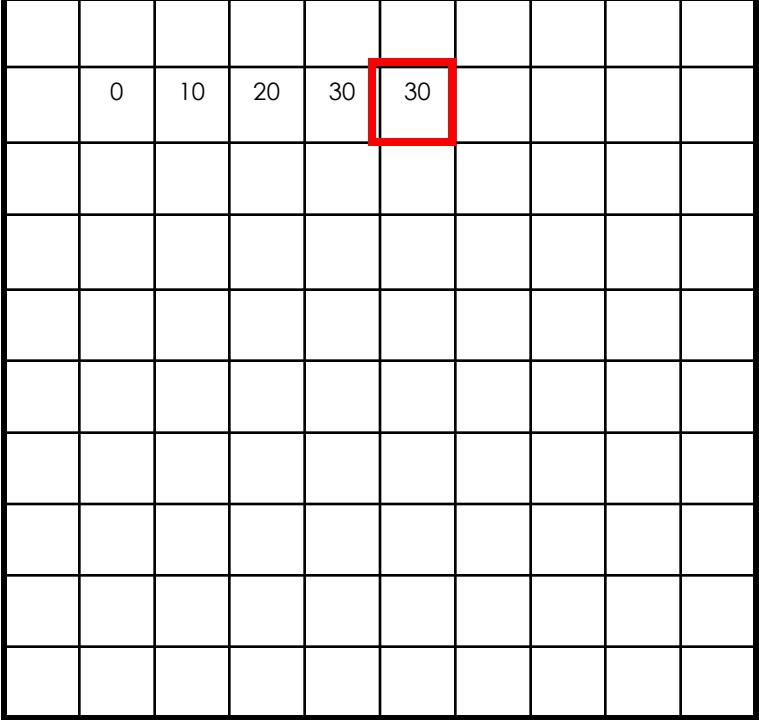
Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

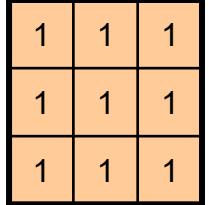


The grid shows the result of applying the filter h to the input image f . The value 30 is highlighted in a red box at position (4,4), indicating the center of the receptive field for that output unit.

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


$$f[., .]$$

$$h[., .]$$

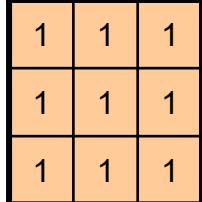
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30				

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

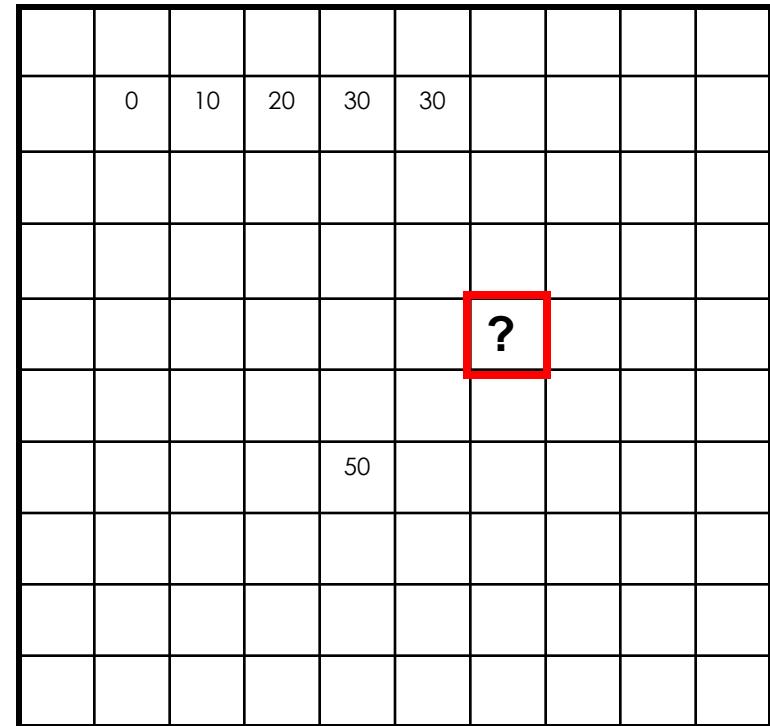
$$g[\cdot, \cdot] \frac{1}{9}$$


A 3x3 matrix where every element is 1/9. The matrix is enclosed in a black border.

$$f[., .]$$

$$h[., .]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$g[\cdot, \cdot] \quad \frac{1}{9} \begin{array}{|ccc|} \hline 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ \hline \end{array}$$

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[\cdot, \cdot]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Box Filter

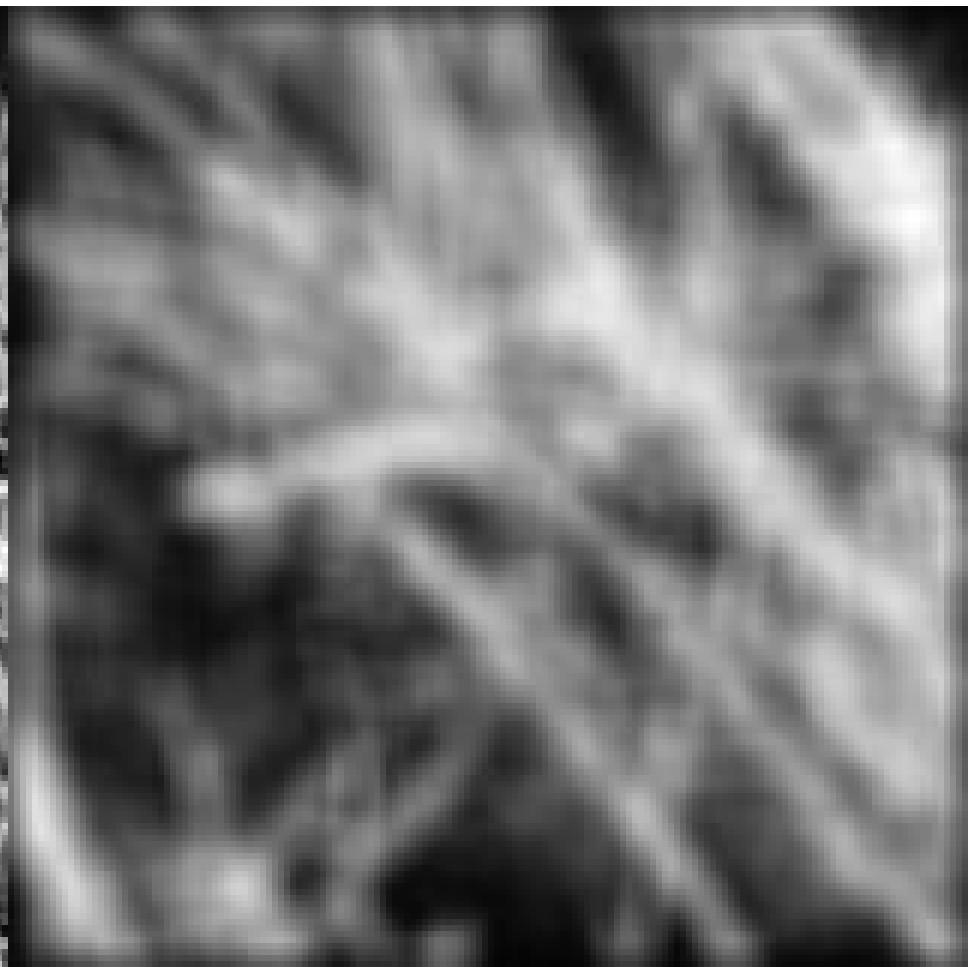
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

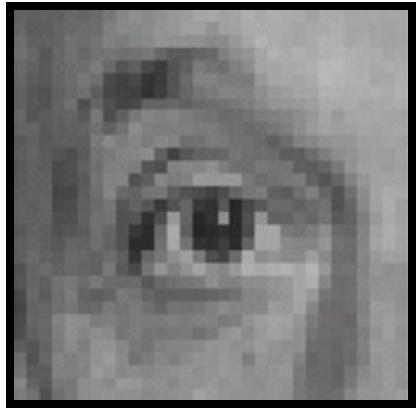
$g[\cdot, \cdot]$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Smoothing with box filter



Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

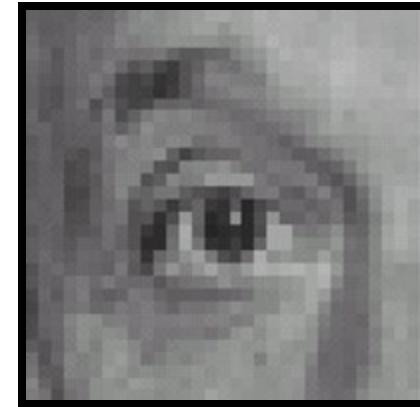
?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

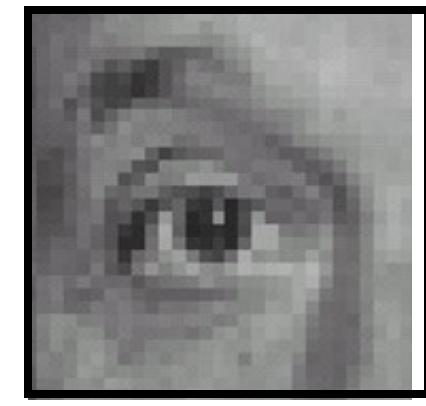
?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Practice with linear filters



0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

Original

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

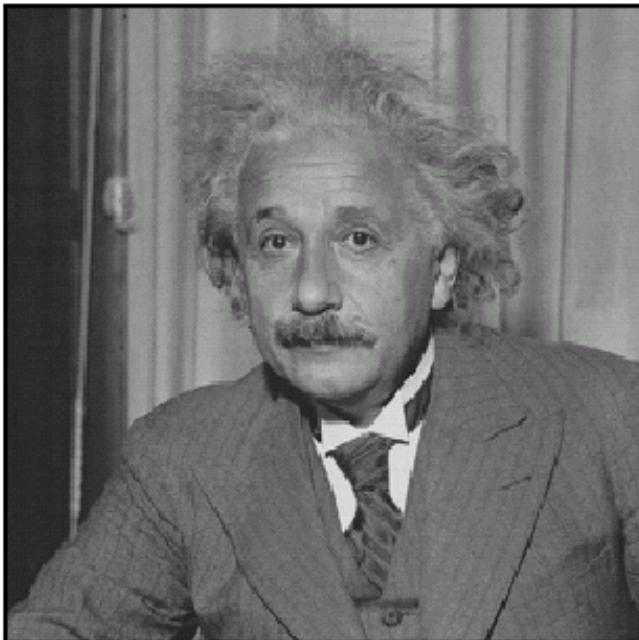
$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1



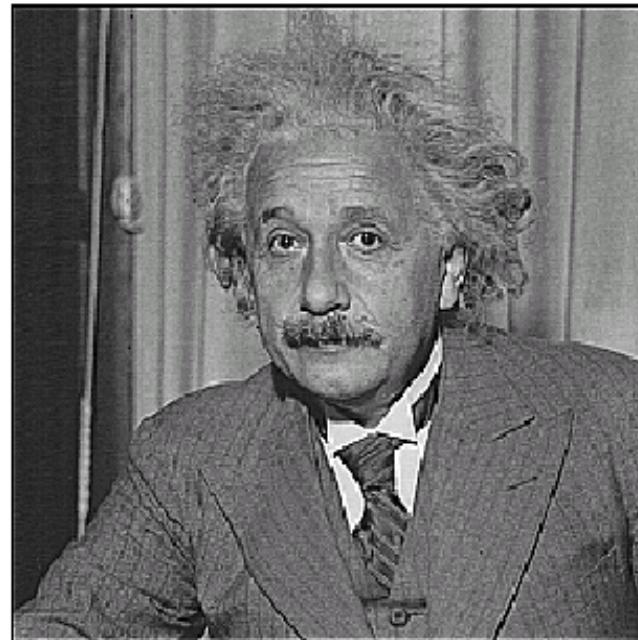
Sharpening filter

- Accentuates differences with local average

Sharpening

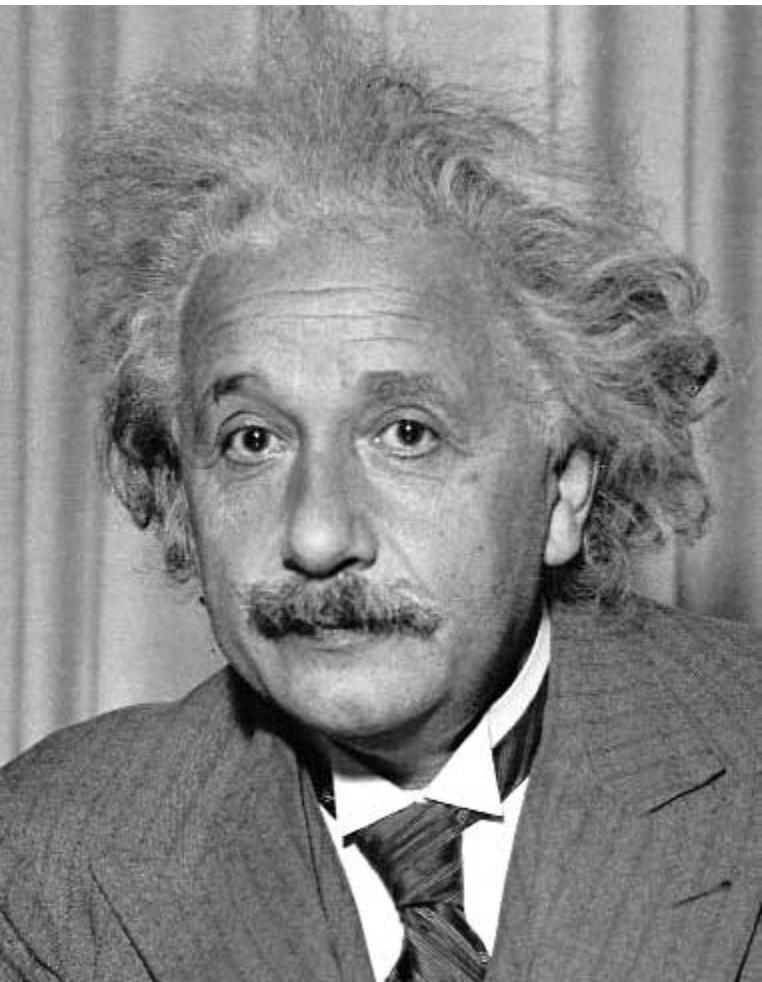


before



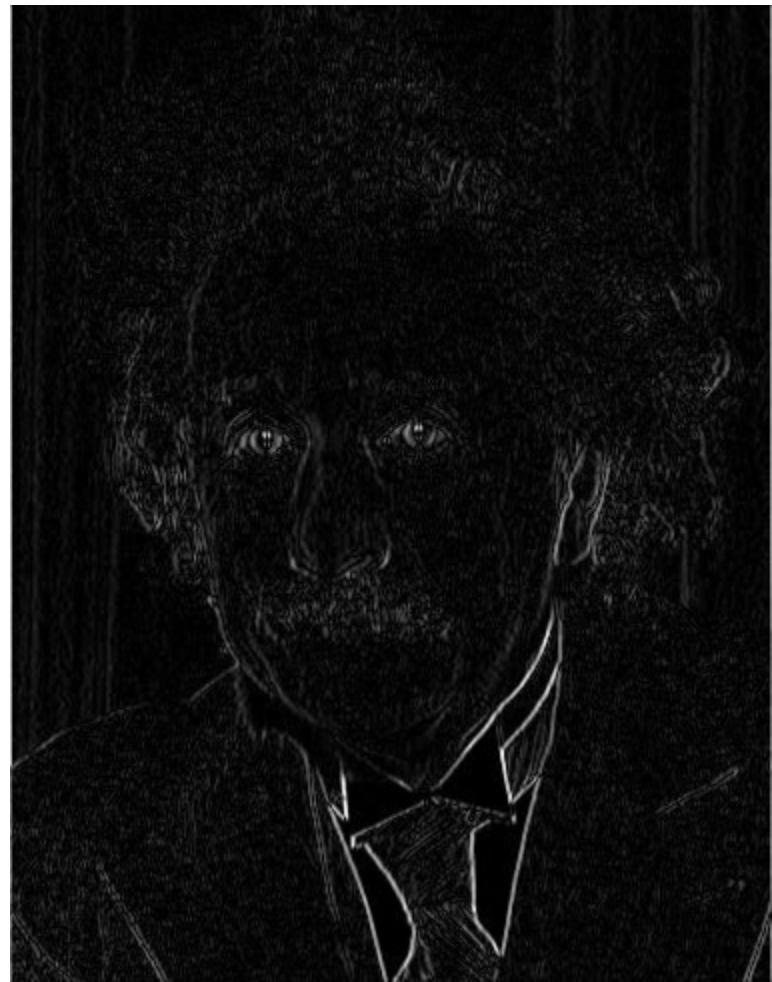
after

Other filters



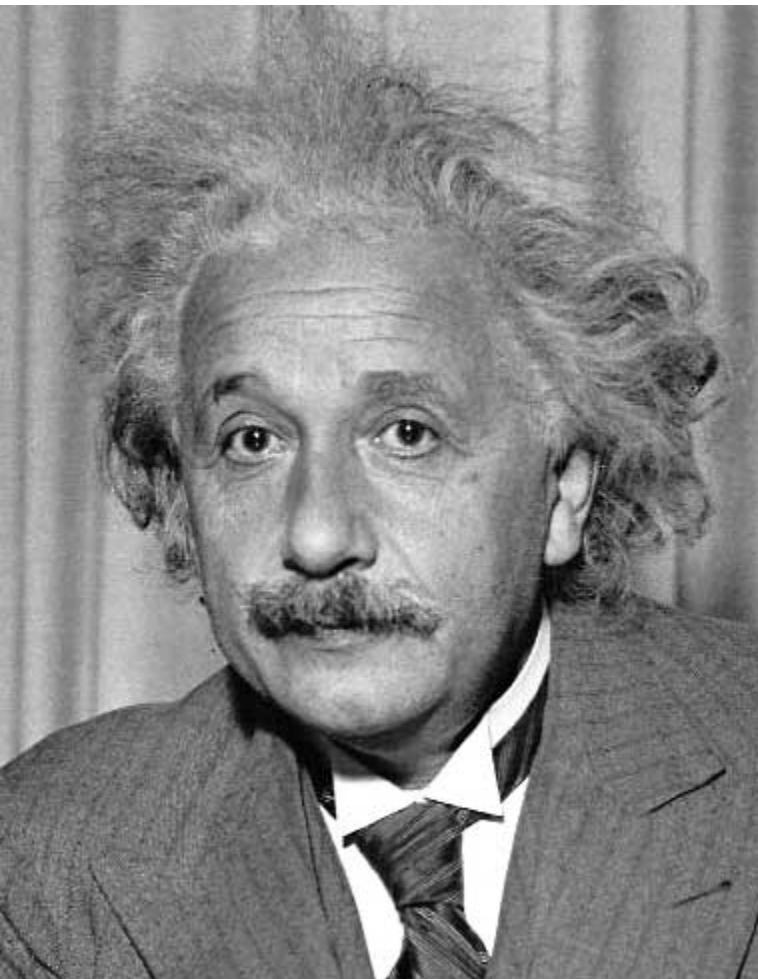
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Other filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

Filtering vs. Convolution

- 2d filtering $g = \text{filter}$ $f = \text{image}$

– $h = \text{filter2}(g, f);$ or
 $h = \text{imfilter}(f, g);$

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

- 2d convolution

– $h = \text{conv2}(g, f);$

$$h[m, n] = \sum_{k, l} g[k, l] f[m - k, n - l]$$

Key properties of linear filters

Linearity:

$$\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$$

Shift invariance: same behavior regardless of pixel location

$$\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$$

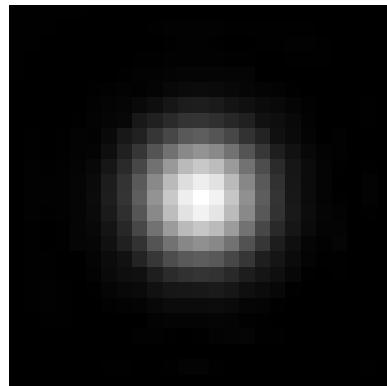
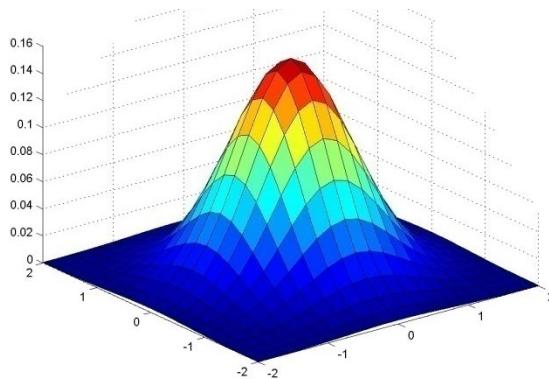
Any linear, shift-invariant operator can be represented as a convolution

More properties

- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
 - But particular filtering implementations might break this equality
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$
- Identity: unit impulse $e = [0, 0, 1, 0, 0]$,
 $a * e = a$

Important filter: Gaussian

- Weight contributions of neighboring pixels by nearness

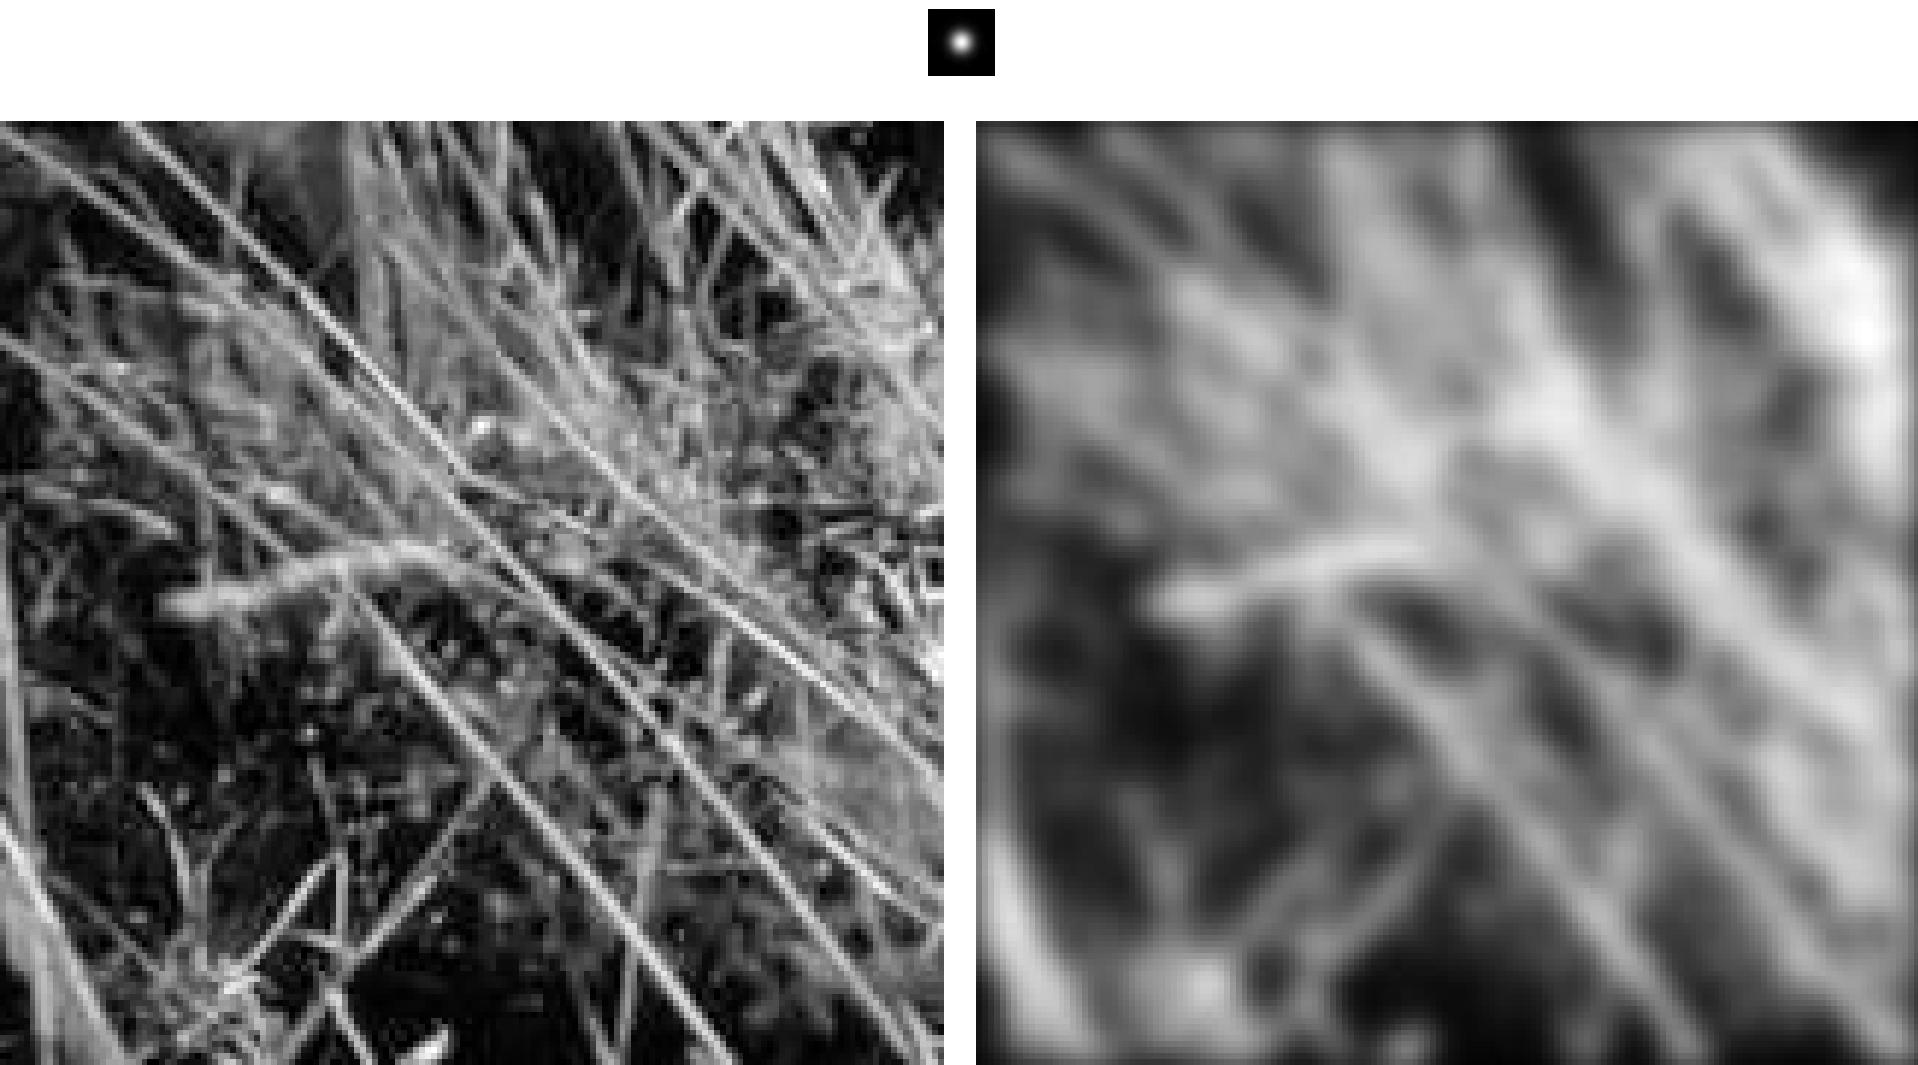


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

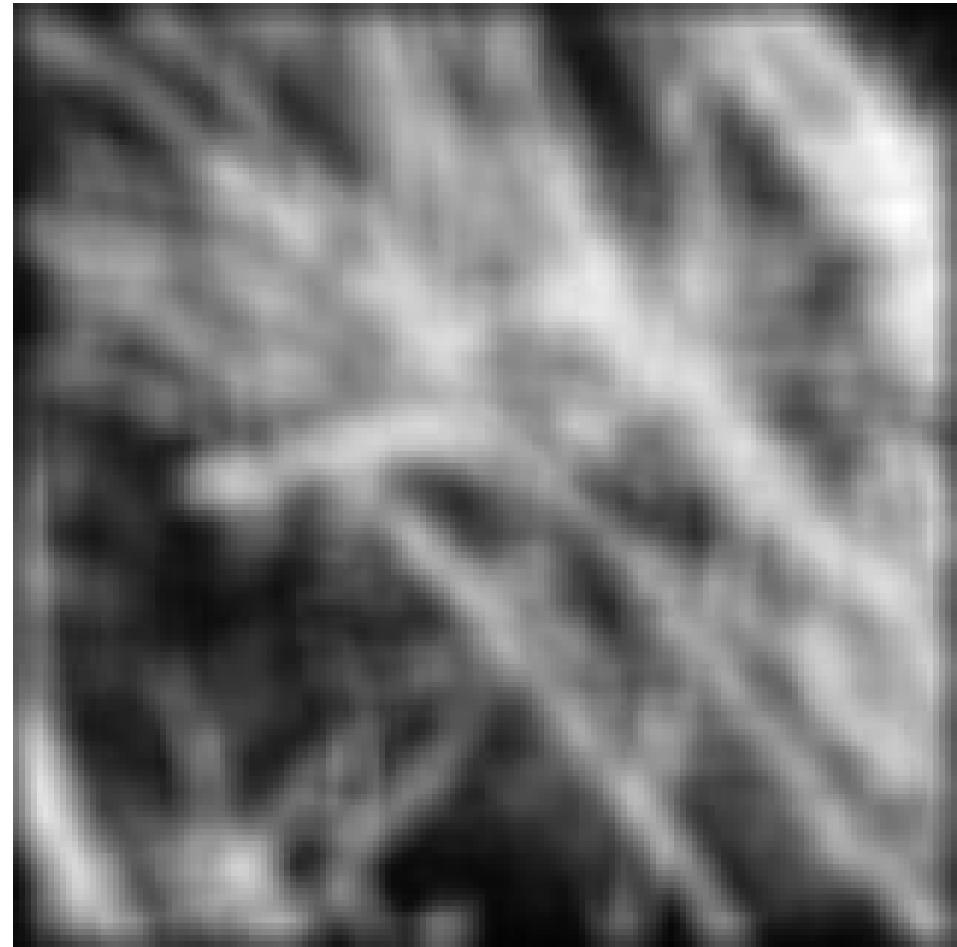
$5 \times 5, \sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Smoothing with Gaussian filter



Smoothing with box filter



Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
 - Images become more smooth
- Convolution with self is another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians

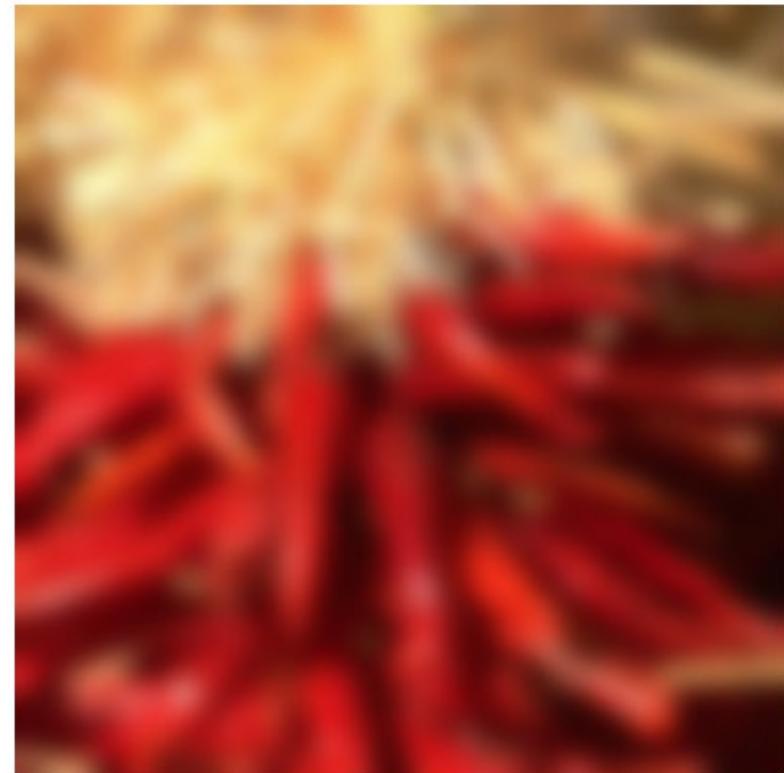
Practical matters

How big should the filter be?

- Values at edges should be near zero
- Rule of thumb for Gaussian: set filter half-width (h) to about 3σ

Practical matters

- What about near the edge?
 - the filter window falls off the edge of the image
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge

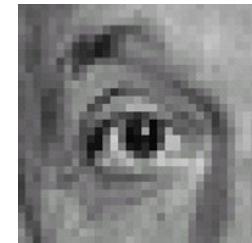


Practical matters

- methods (MATLAB):
 - clip filter (black): `imfilter(f, g, 0)`
 - wrap around: `imfilter(f, g, 'circular')`
 - copy edge: `imfilter(f, g, 'replicate')`
 - reflect across edge: `imfilter(f, g, 'symmetric')`

Take-home messages

- Linear filtering is sum of dot product at each position
 - Can smooth, sharpen, translate (among many other uses)
- Be aware of details for filter size, extrapolation, cropping



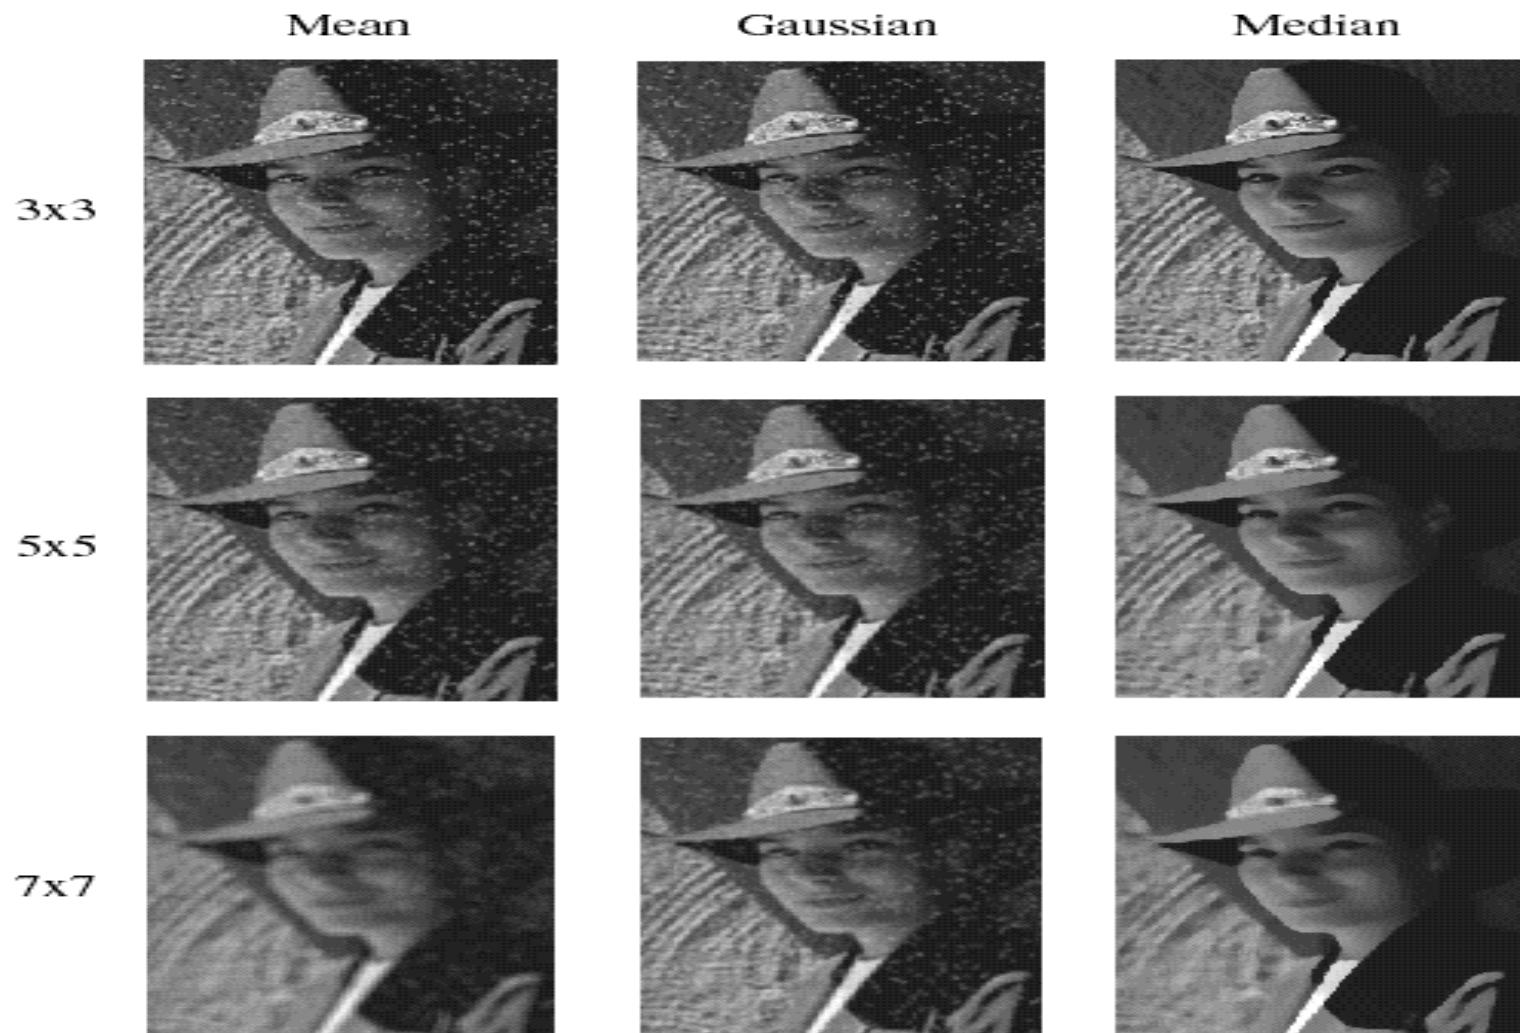
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$
A 3x3 kernel matrix representing a linear filter. Each element is labeled with the value '1'. The matrix is multiplied by $\frac{1}{9}$ to represent a uniform average filter.



Median

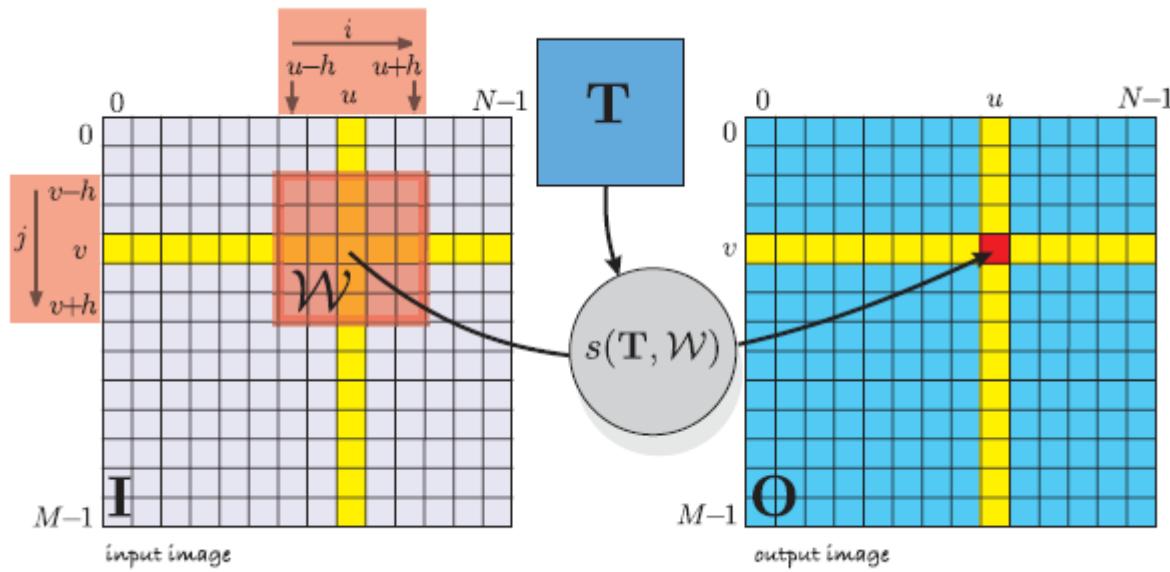
- A **Median Filter** operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?

Comparison: salt and pepper noise



TEMPLATE MATCHING

- Gaussian Kernel as a matrix but also an image
- Template are images used as kernels.
- Need to identify parts of an image that are most similar to your template

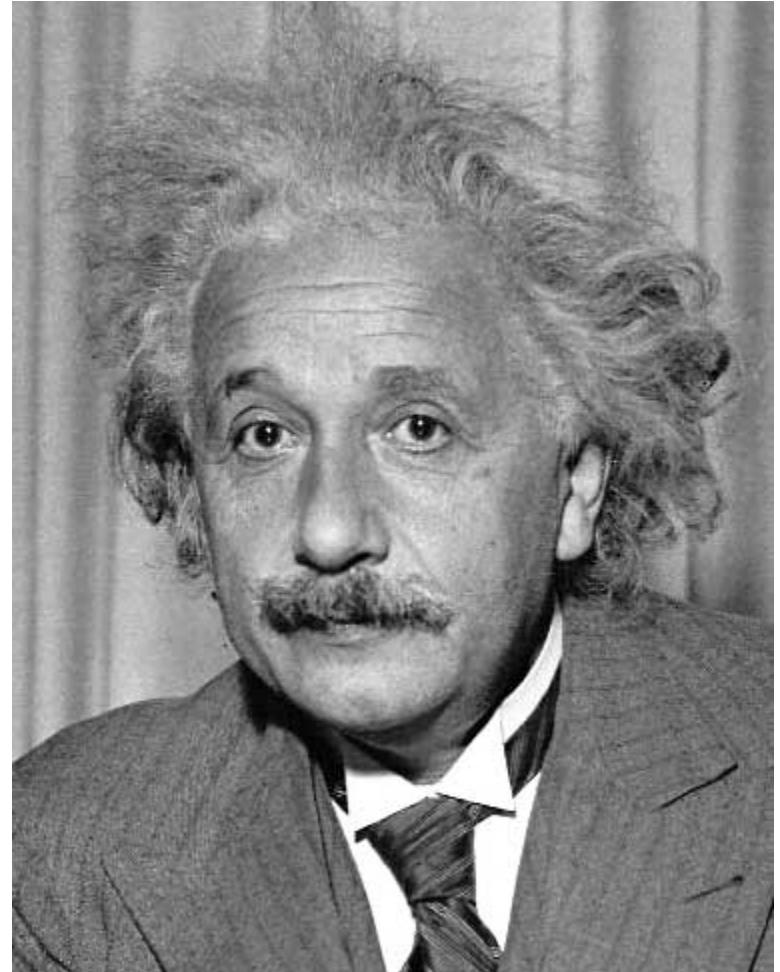


$$O[u, v] = s(T, W), \forall (u, v) \in I$$

Where T is the $w \times w$ template and W is the $w \times w$ window centered at (u, v)

Template matching

- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches?
 - Correlation
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation



Multiple similarity measures

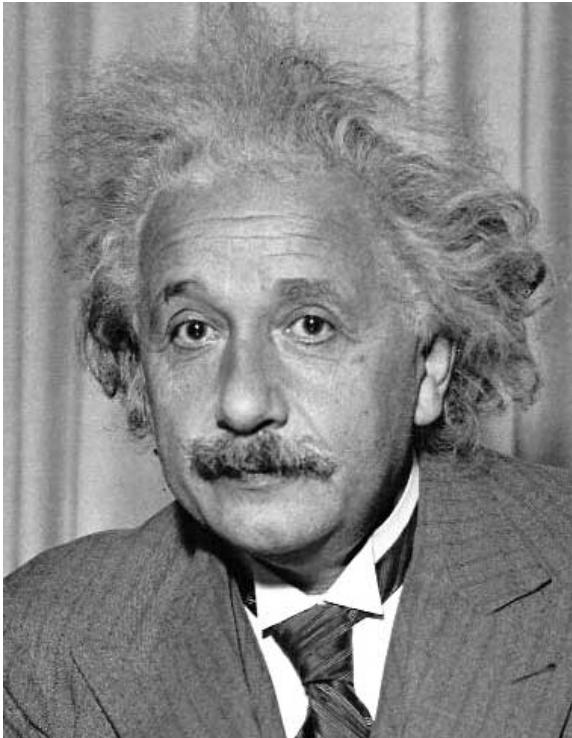
- Sum of the absolute differences (SAD)
 - $s = \sum_{(u,v) \in I} |I_1[u, v] - I_2[u, v]|$
- Sum of the squared differences (SSD)
 - $s = \sum_{(u,v) \in I} (I_1[u, v] - I_2[u, v])^2$
- Zero mean SAD (ZSAD)
 - $s = \sum_{(u,v) \in I} |(I_1[u, v] - \bar{I}_1) - (I_2[u, v] - \bar{I}_2)|$
- Zero mean SSD (ZSSD)
 - $s = \sum_{(u,v) \in I} ((I_1[u, v] - \bar{I}_1) - (I_2[u, v] - \bar{I}_2))^2$
- Normalized cross-correlation (NCC)
 - $s = \frac{\sum_{(u,v) \in I} I_1[u, v] \cdot I_2[u, v]}{\sqrt{\sum_{(u,v) \in I} I_1^2[u, v] \cdot \sum_{(u,v) \in I} I_2^2[u, v]}}$
- zncc subtract the mean first

Matching with filters

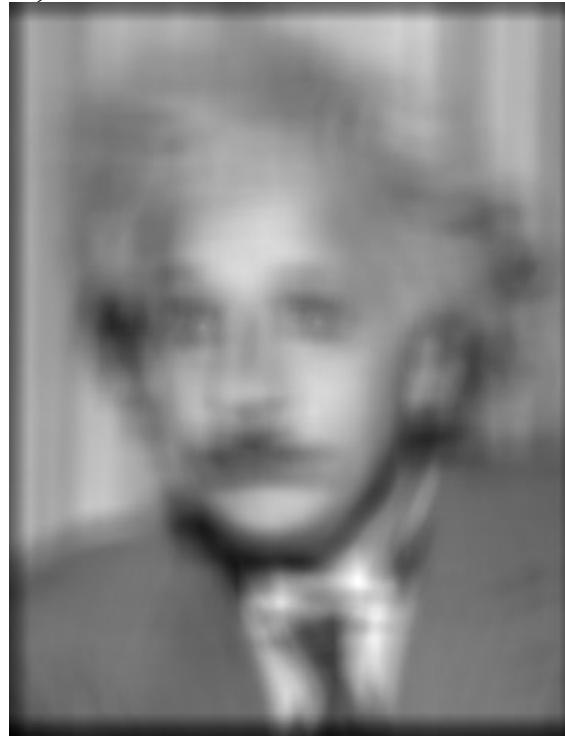
- Goal: find  in image
- Method 0: filter the image with eye patch

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

f = image
g = filter



Input



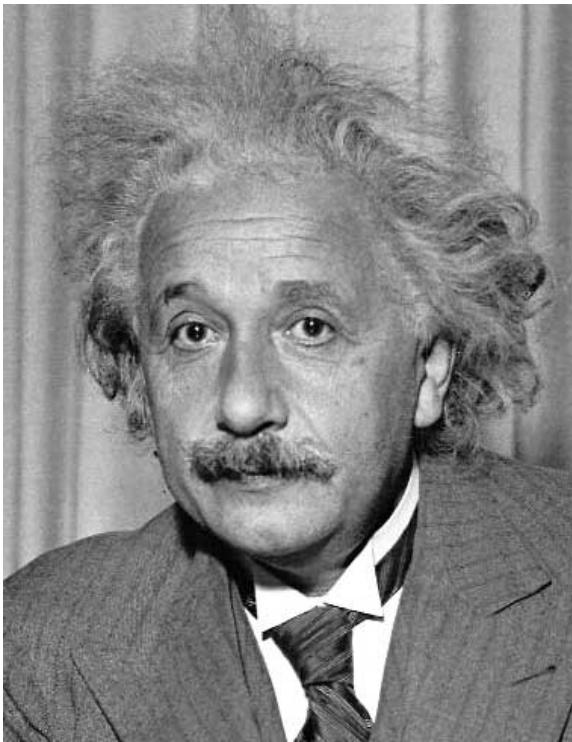
Filtered Image

What went wrong?

Matching with filters

- Goal: find  in image
- Method 1: filter the image with zero-mean eye

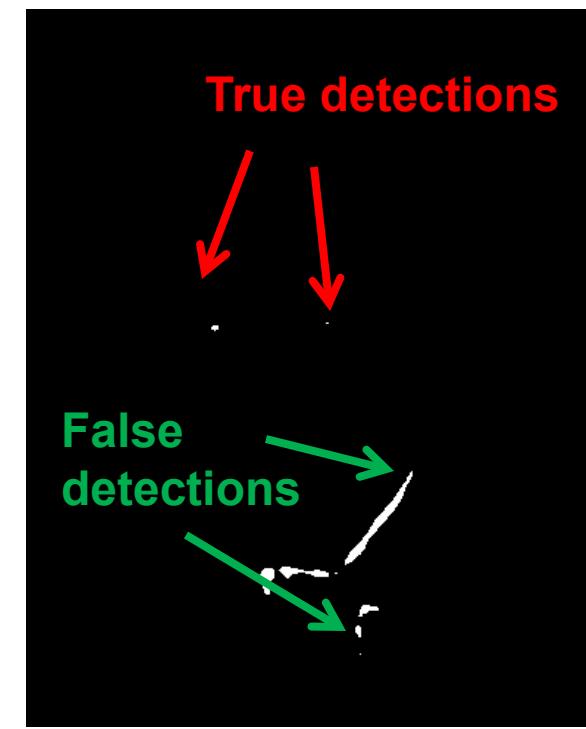
$$h[m, n] = \sum_{k, l} (g[k, l] - \bar{g}) (f[m + k, n + l])$$



Input



Filtered Image (scaled)

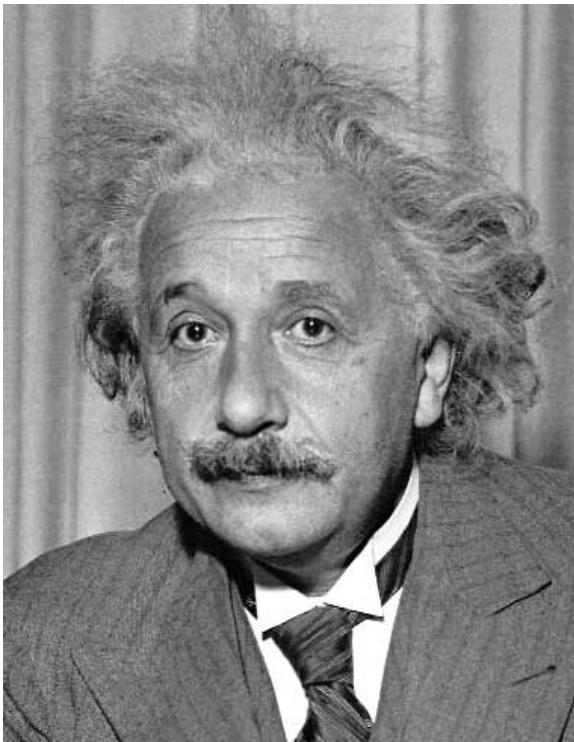


Thresholded Image

Matching with filters

- Goal: find  in image
- Method 2: SSD

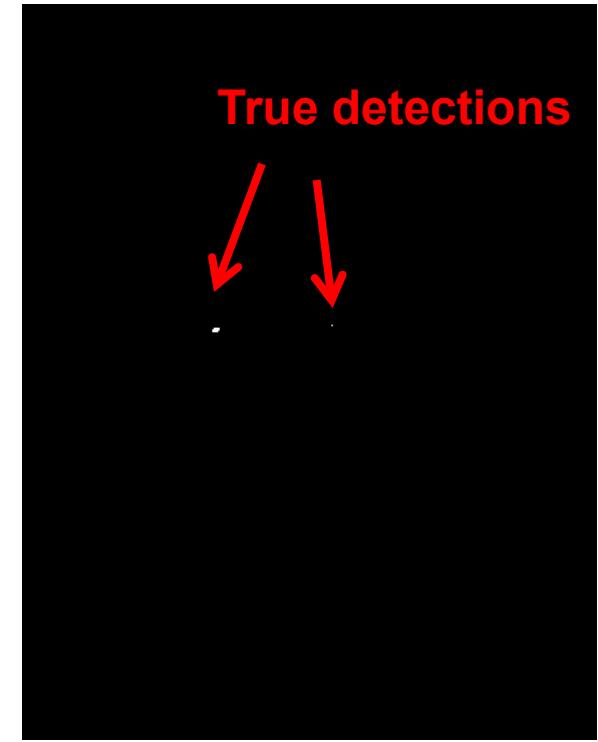
$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1 - $\text{sqrt}(\text{SSD})$



Thresholded Image

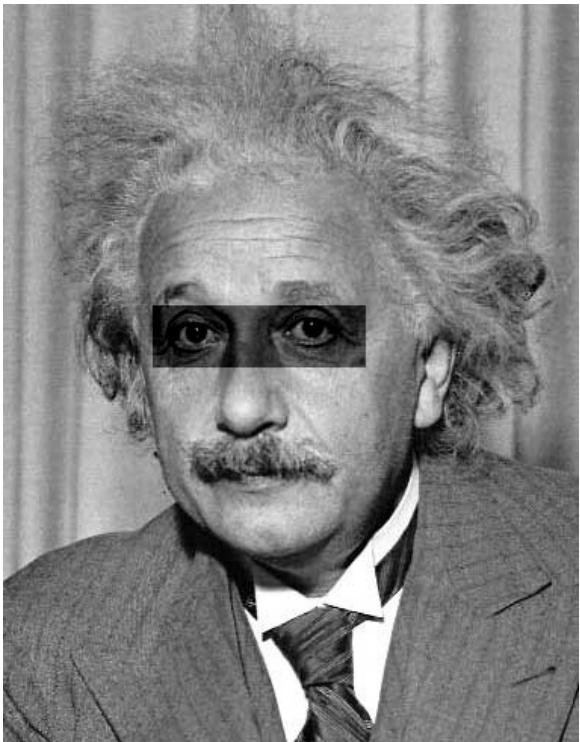
True detections

Matching with filters

- Goal: find  in image
- Method 2: SSD

What's the potential downside of SSD?

$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1- sqrt(SSD)

Matching with filters

- Goal: find  in image
- Method 3: Normalized cross-correlation

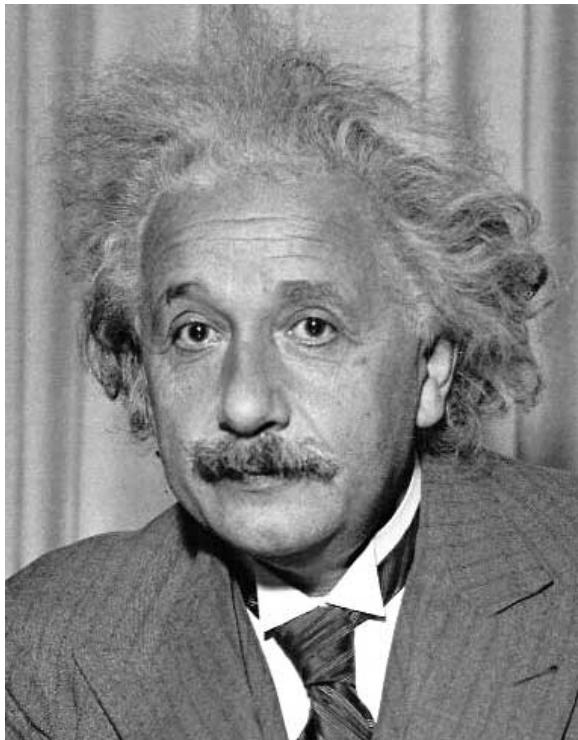
$$h[m, n] = \frac{\sum_{k,l} (g[k, l] - \bar{g})(f[m - k, n - l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k, l] - \bar{g})^2 \sum_{k,l} (f[m - k, n - l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

mean template mean image patch
↓ ↓

Matlab: `normxcorr2(template, im)`

Matching with filters

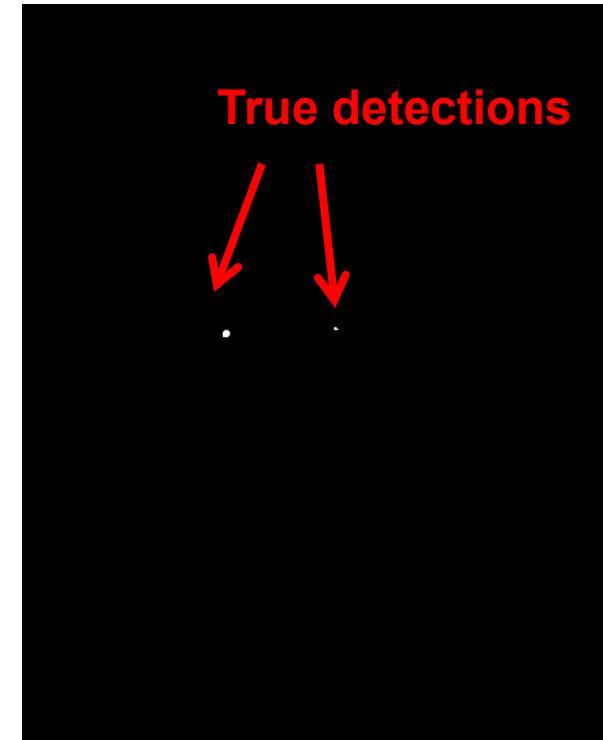
- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



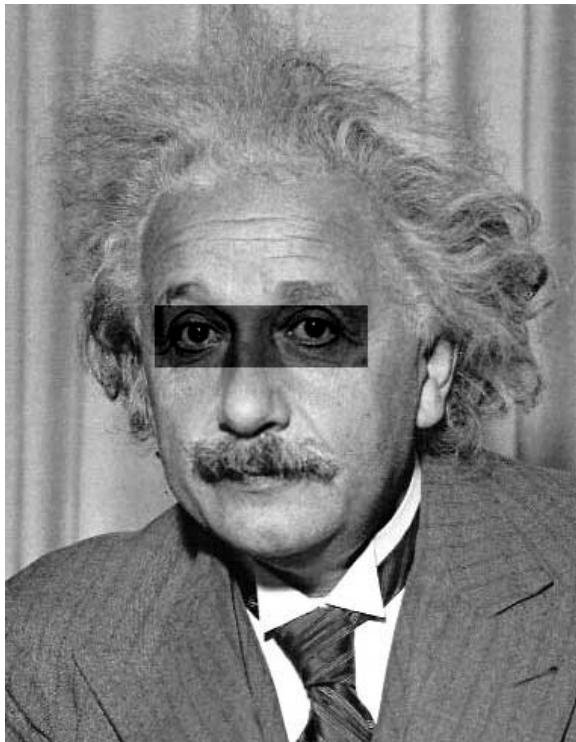
Normalized X-Correlation



True detections
Thresholded Image

Matching with filters

- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



Thresholded Image

True detections

Practical issues

- Not invariant to scale, max 10-20% change in scale between T and W.
- Mixed pixel problem where background is also part of the image. As image moves, background changes... lowering similarity.
- Center of the image is different than edges due to distortion.
- False matching can occur. Use heuristics to reduce that

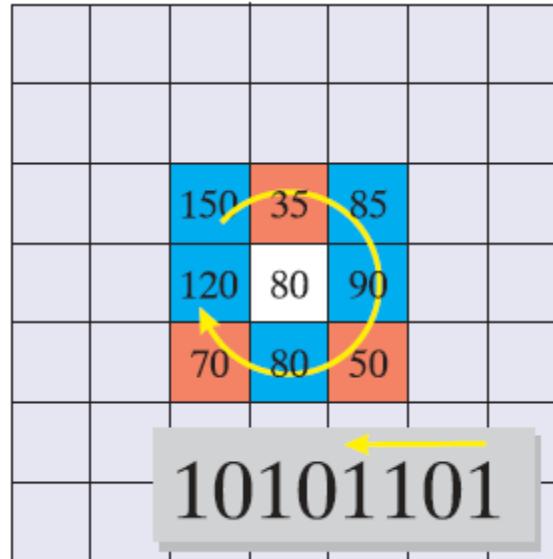
Q: What is the best method to use?

A: Depends

- SSD: faster, sensitive to overall intensity
- Normalized cross-correlation: slower, invariant to local average intensity and contrast
- But really, neither of these baselines are representative of modern recognition.

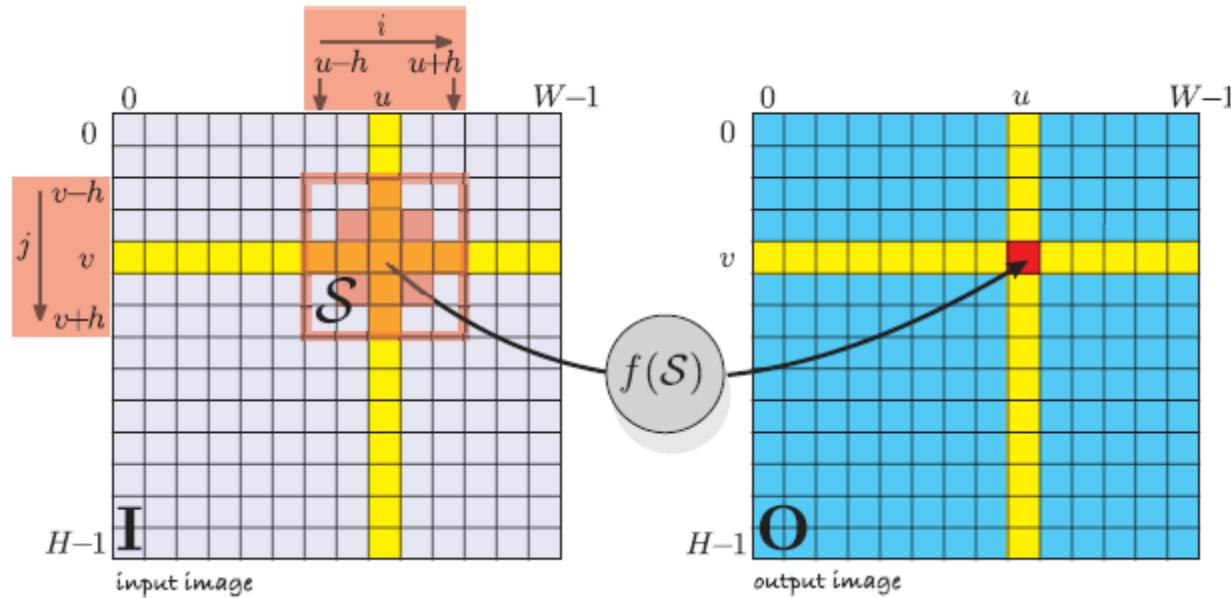
Non-parametric similarity

Non-parametric similarity are more robust to mixed pixel like census and rank
Compares pixels to the center pixel and evaluate using hamming distance

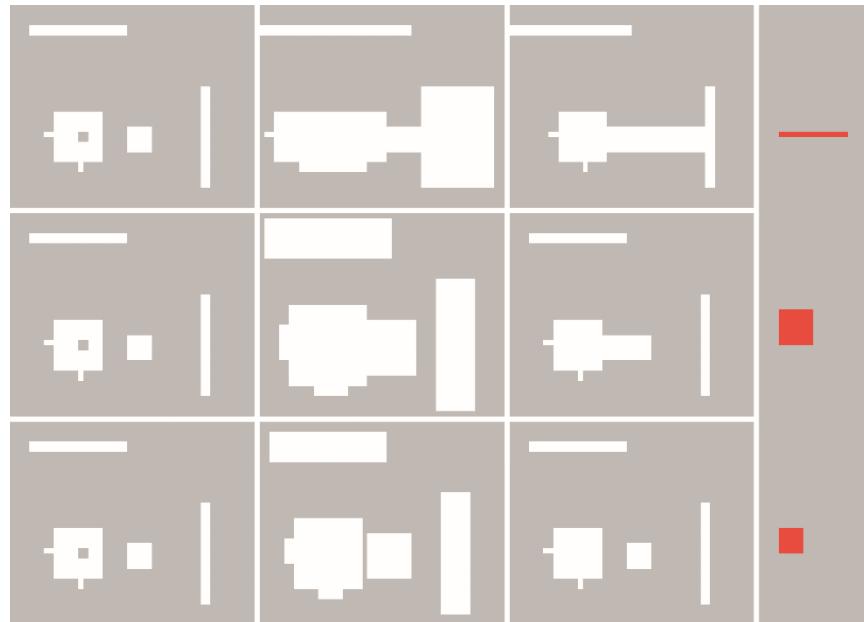


Morphology operations

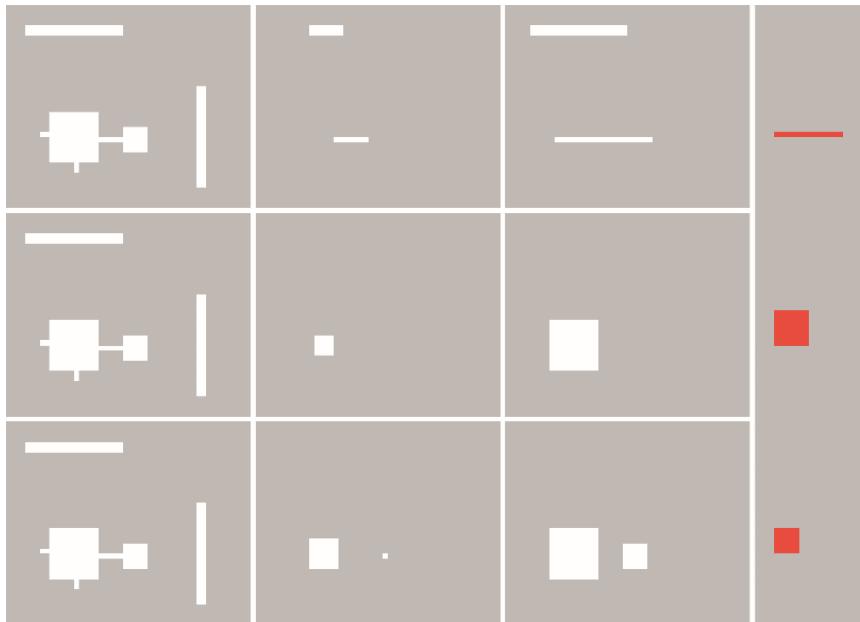
- Non-linear spatial operators used on binary images
 - $O[u, v] = f(I[u + i, v + j]), \forall (i, j) \in S, \forall (u, v) \in I$



Close operation – dilation then erosion



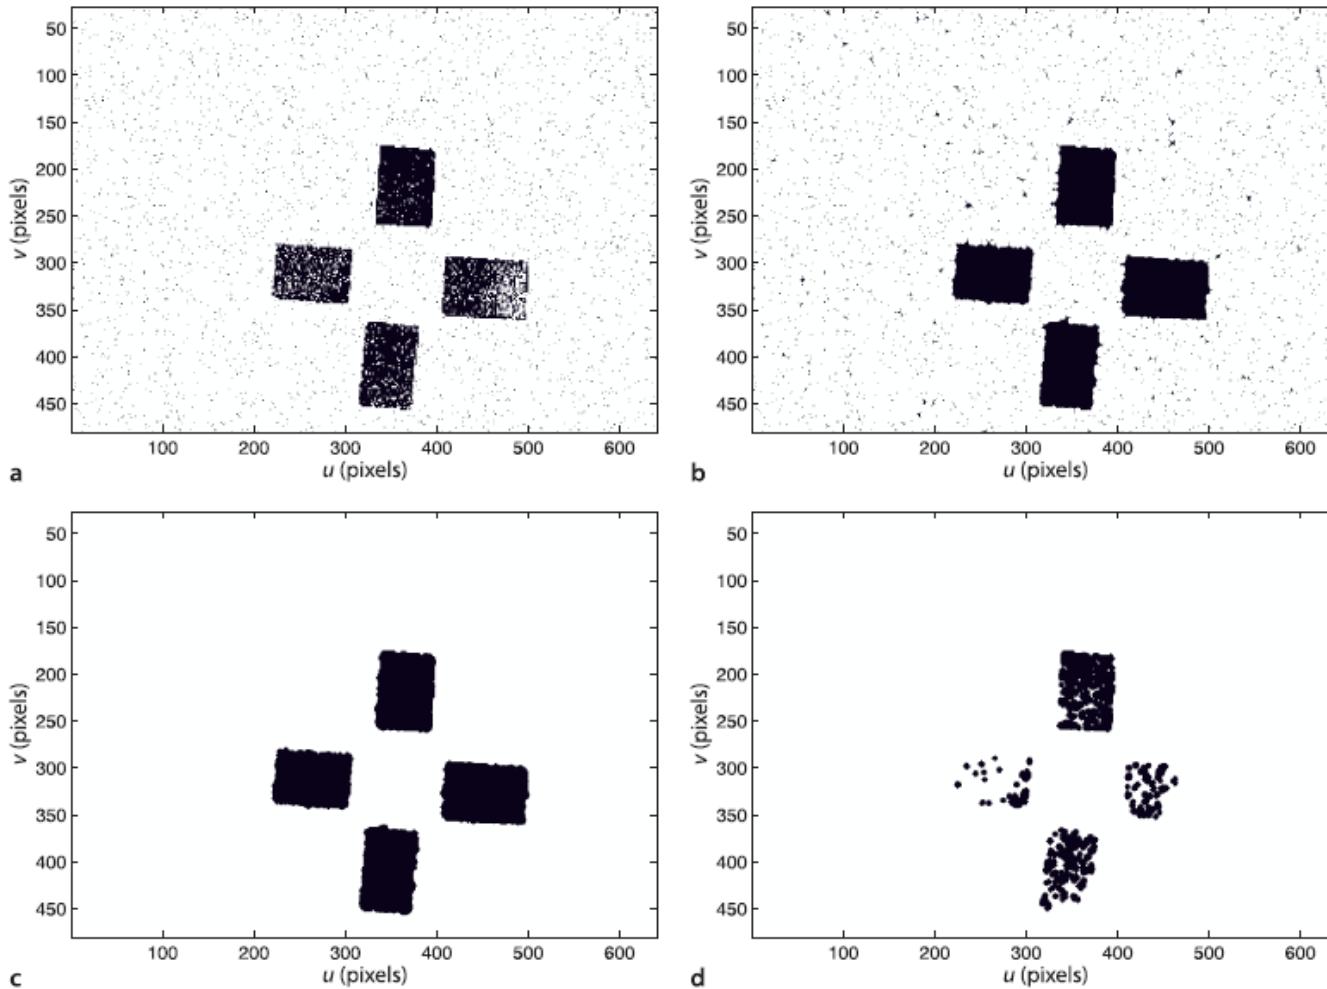
Open operation – erosion then dilation



Matlab

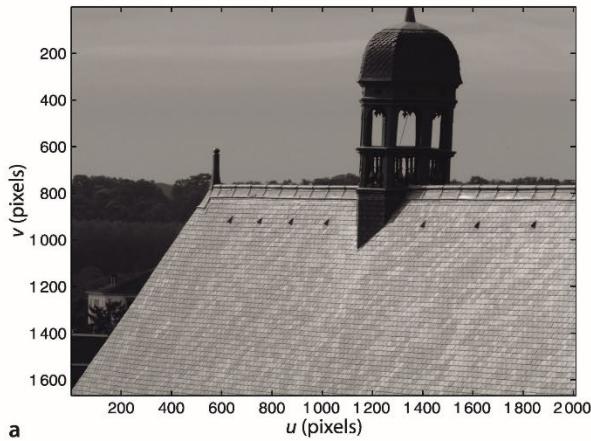
- `imorph(IM, SE, OP)` is the image IM after morphological processing with the operator OP and structuring element SE.
- Or use
 - `ierode (im, S)`
 - `Idilate (im, S)`
- `iopen`: erosion then dilation
- `iclose`: dilation then erosion

Example of noise removal

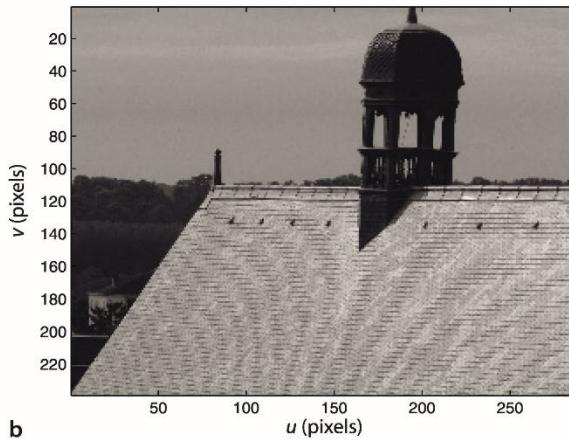


a: original, b: close, c: open, d: the opposite sequence

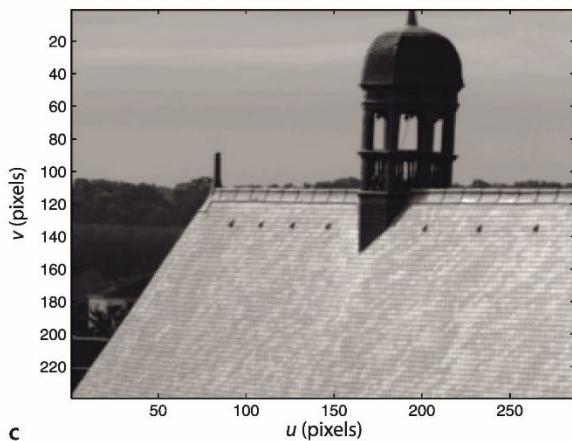
Image resizing



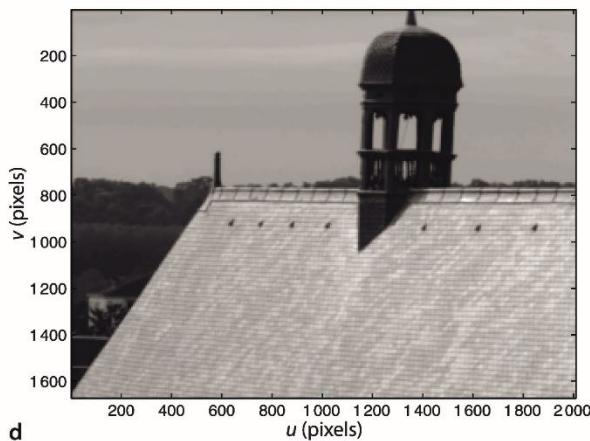
a



b



c



d

Image pyramids

- Pyramidal decomposition of an image
- Used for coarse-to-fine strategies
- Use ipyramid



FEATURE EXTRACTION

- Image processing: image -> image
- Feature extraction (segmentation): image -> image features or objects
- Information concentration or reduction step
 - Regions
 - Lines
 - Interest or keypoint

Region features

- Segment or extract regions that represent objects
- Three steps
 - Classify: assign pixels to classes (binary or more)
 - Represent: connect regions to form spatial labeled sets
 - Describe: size, position, shape, etc.

Binary classification

$$c[u, v] = \begin{cases} 0 & I[u, v] < t \\ 1 & I[u, v] \geq t \end{cases} \quad \forall (u, v) \in I$$

- Matlab example
 - Basic histogram thresholding
 - More advanced thresholding
 - N. Otsu, A Threshold Selection Method from Gray-Level Histograms, IEEE Trans. Systems, Man and Cybernetics, Vol SMC-9(1), Jan 1979, pp 62-66
 - W. Niblack, An Introduction to Digital Image Processing, Prentice-Hall, 1986.
 - J. Matas, O. Chum, M. Urban, and T. Pajdla, Robust wide-baseline stereo from maximally stable extremal regions", Image and Vision Computing, vol. 22, pp. 761-767, Sept. 2004.

Maximally Stable Extremal Regions

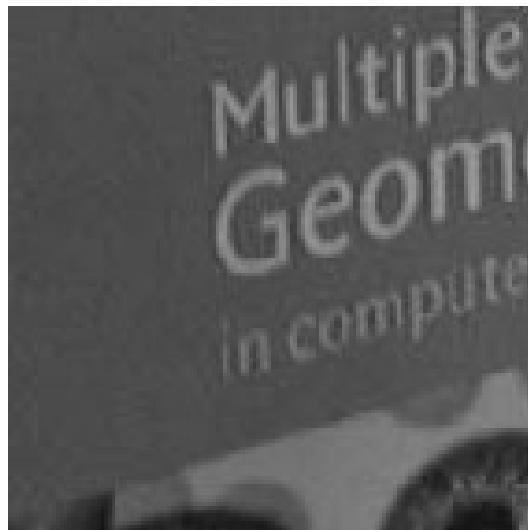
J. Matas, O. Chum, M. Urban, and T. Pajdla, Robust wide-baseline stereo from maximally stable extremal regions", Image and Vision Computing, vol. 22, pp. 761-767, Sept. 2004.

- Maximally Stable Extremal Regions
 - *Threshold* image intensities: $I > \text{thresh}$ for several increasing values of thresh
 - Extract *connected components* ("Extremal Regions")
 - Find a threshold when region is "Maximally Stable", i.e. *local minimum* of the relative growth
 - Approximate each region with an *ellipse*
 - *Affine invariant*



Maximally Stable Extremal Regions [Matas '02]

- Based on Watershed segmentation algorithm
- Select regions that stay stable over a large parameter range



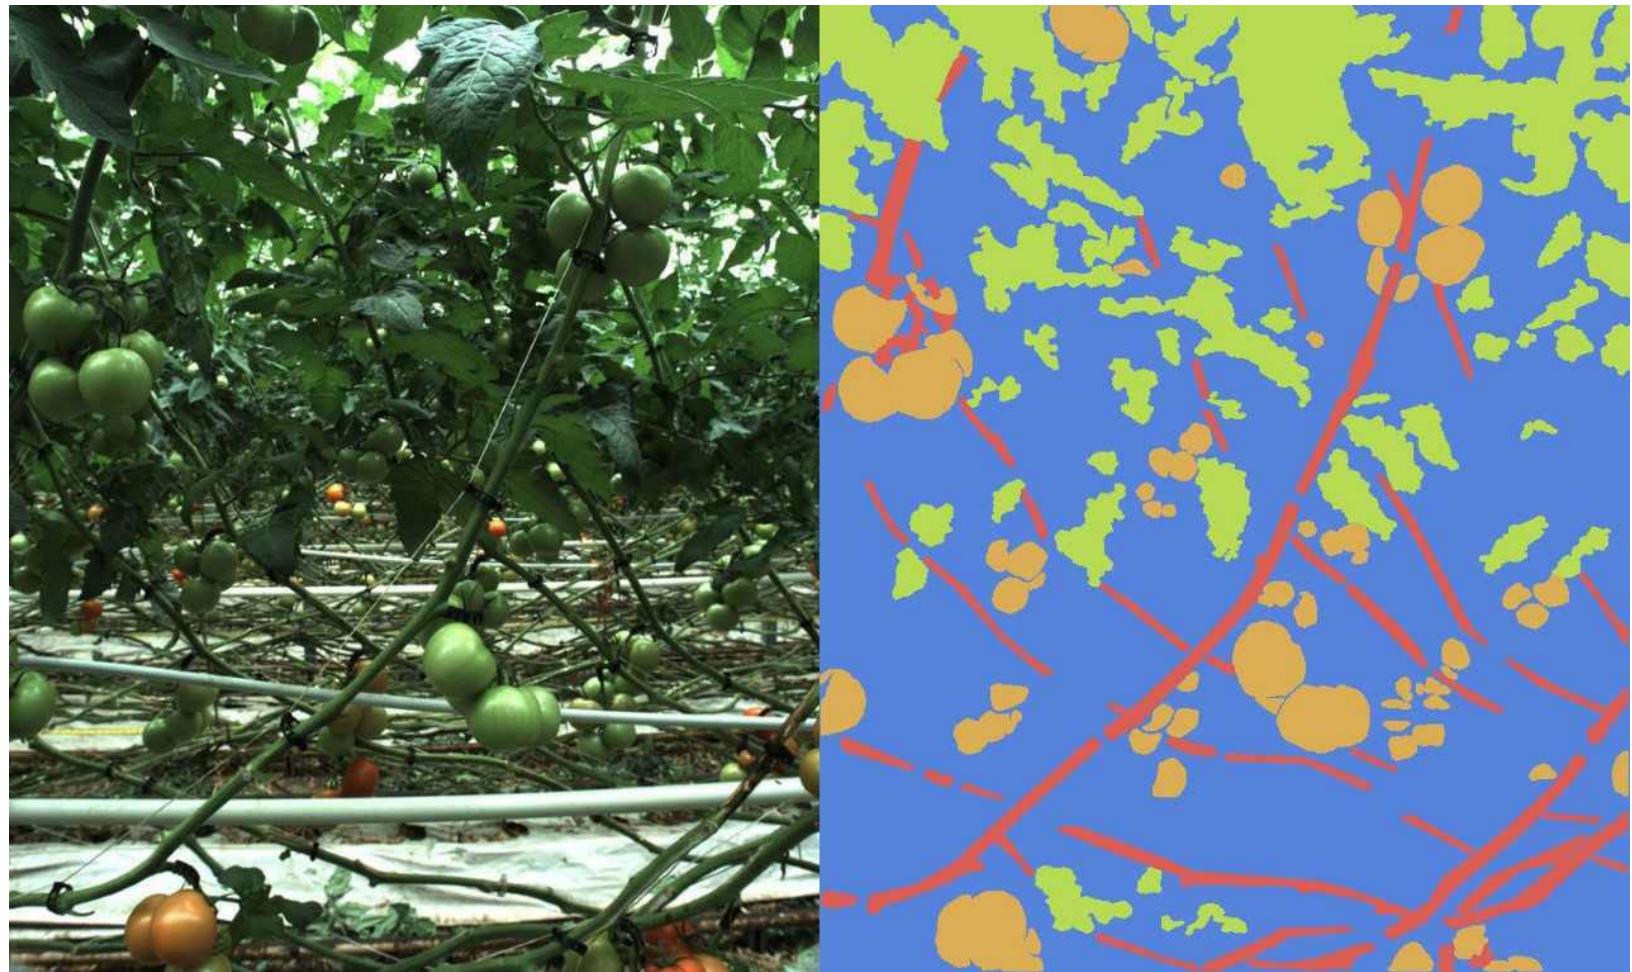
Other segmentation

- Many other region based techniques were proposed including using graph based segmentation
 - P. Felzenszwalb and D. Huttenlocher, "Efficient graph-based image segmentation", Int. Journal on Computer Vision, vol. 59, pp. 167-181, Sept. 2004.

Description of image feature

- Bounding box or area
- Moments: describe center of image (center of mass), shape, and orientation
- Aspect ratio of major axes
- Boundary representations, e.g circle

iblobs: calculate useful metrics for each region in an image like size,





Szeliski 4.2, Corke 12

EDGE DETECTION

Edge detection

- **Goal:** Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)

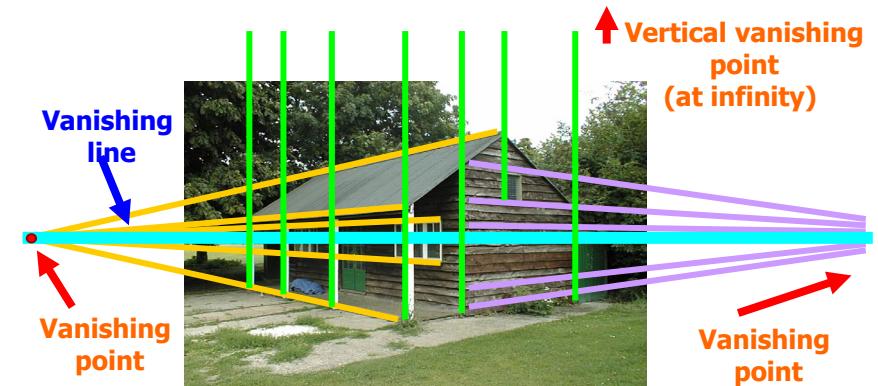


Why do we care about edges?

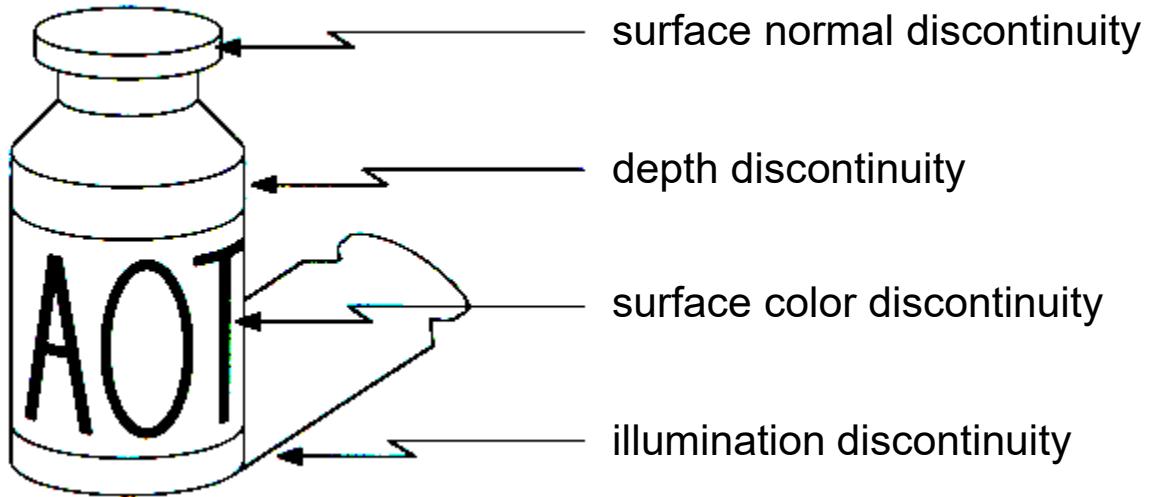
- Extract information,
recognize objects



- Recover geometry and
viewpoint



Origin of Edges



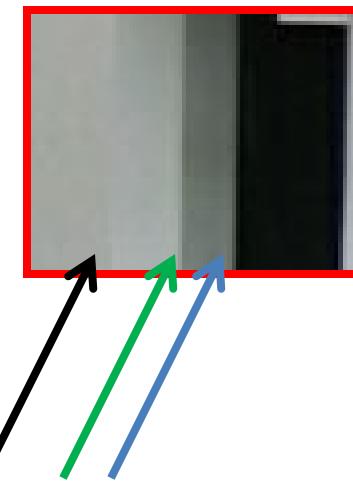
- Edges are caused by a variety of factors

Closeup of edges



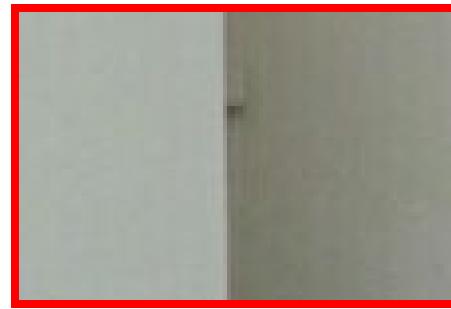
Source: D. Hoiem

Closeup of edges



Source: D. Hoiem

Closeup of edges



Source: D. Hoiem

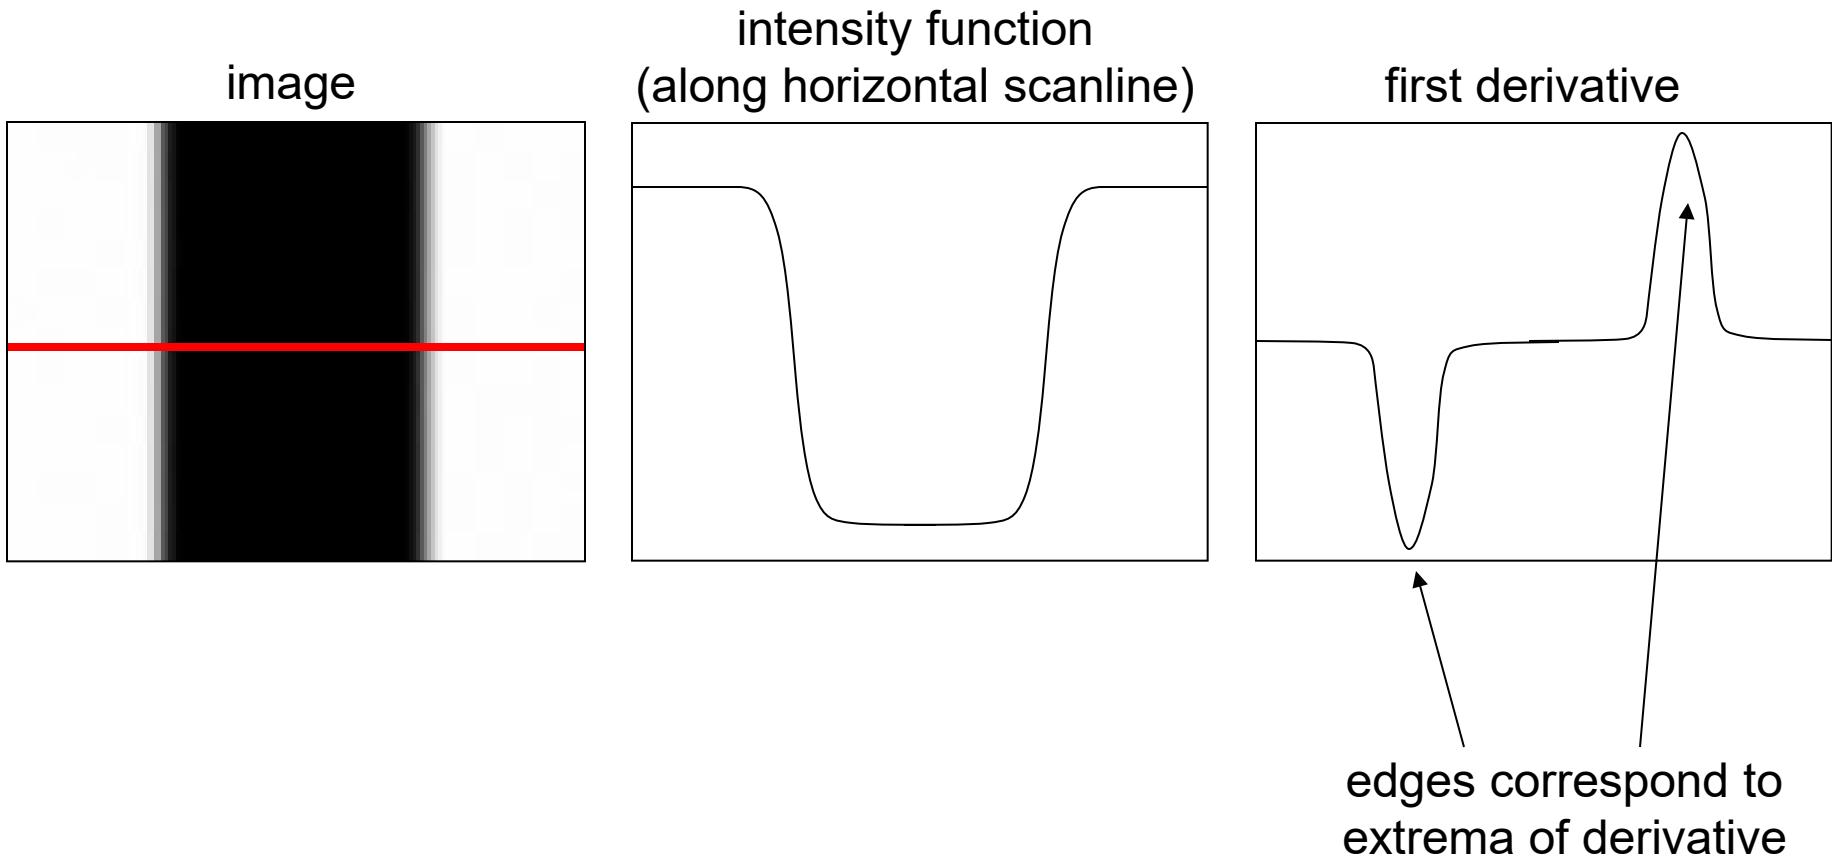
Closeup of edges



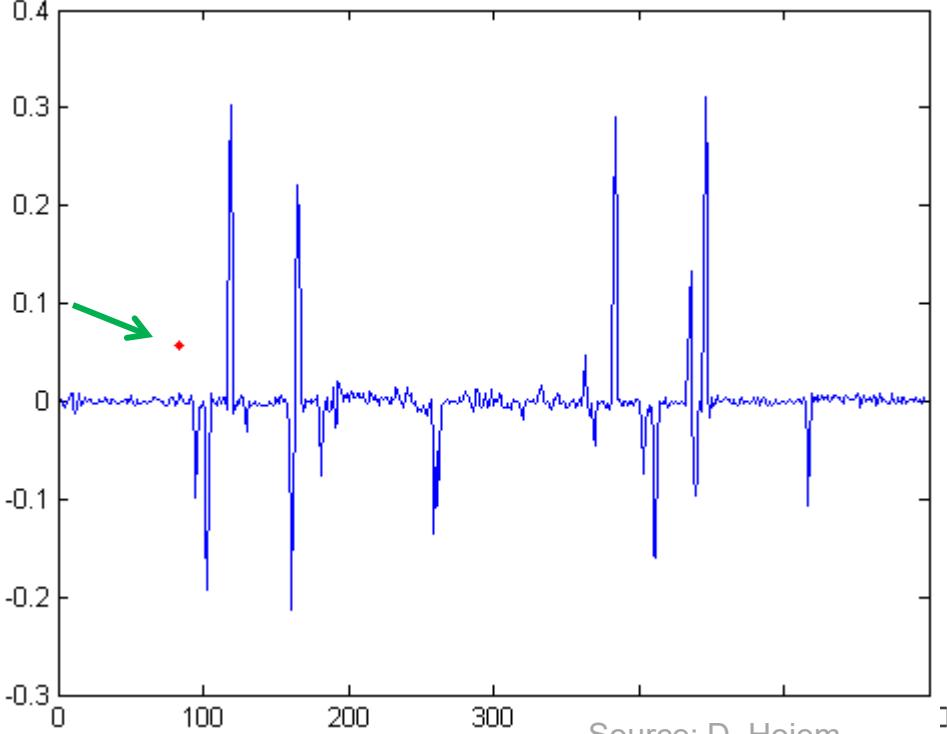
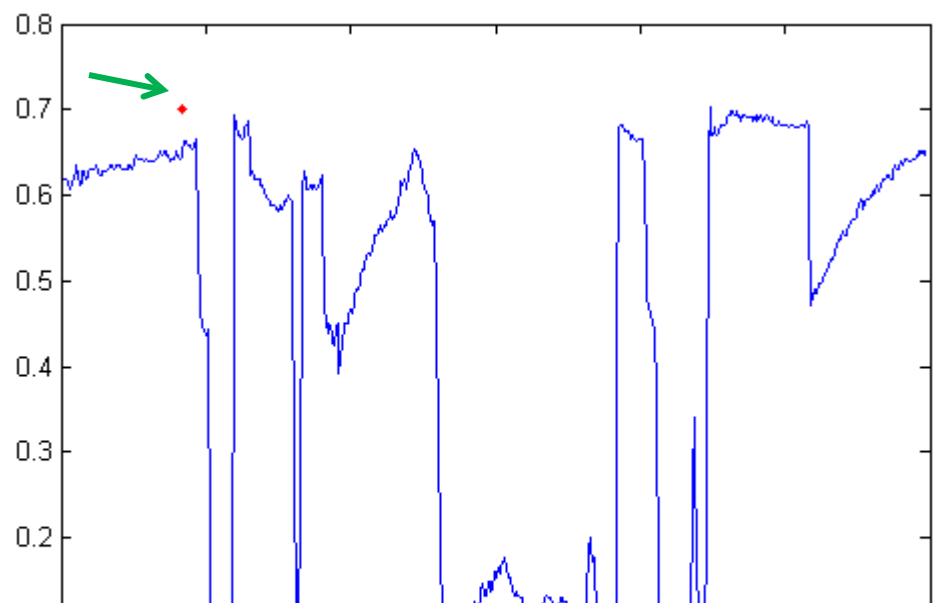
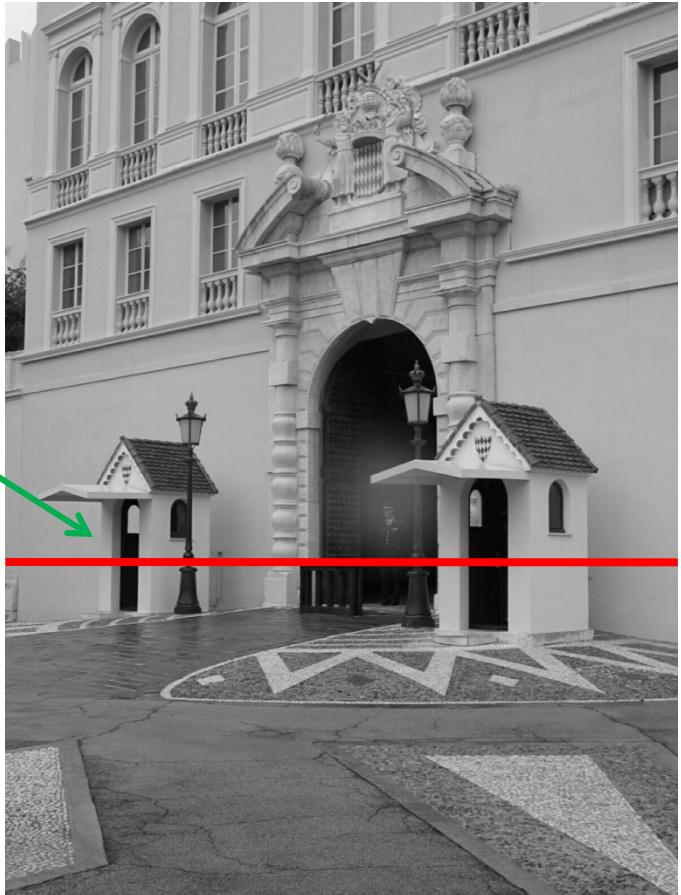
Source: D. Hoiem

Characterizing edges

- An edge is a place of rapid change in the image intensity function

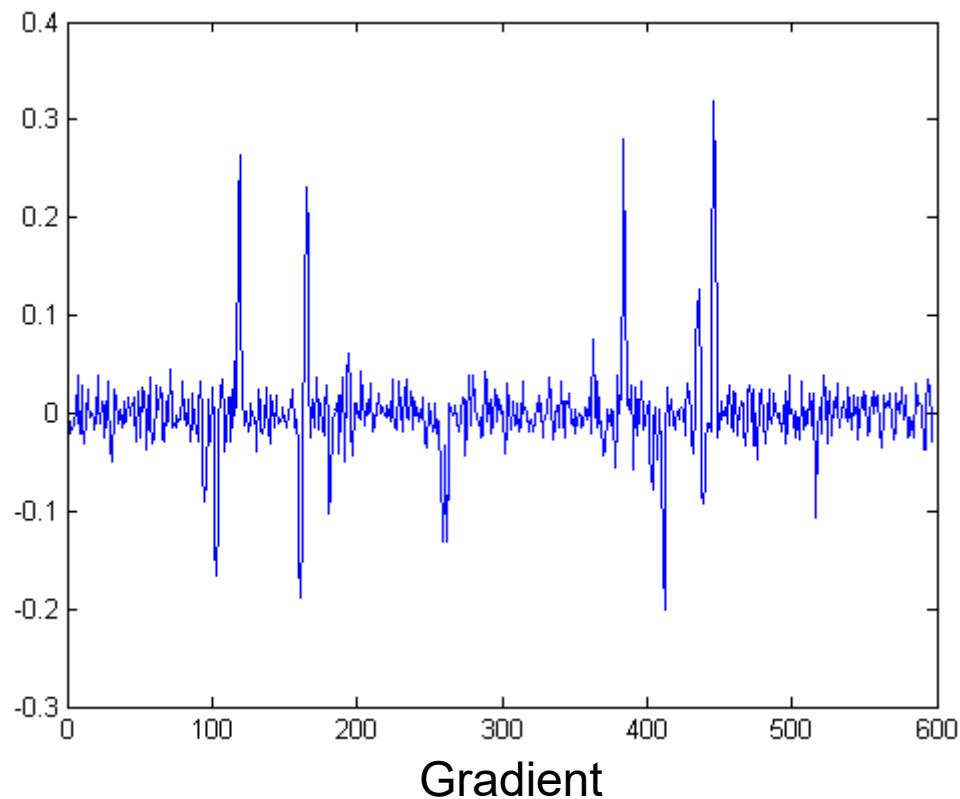
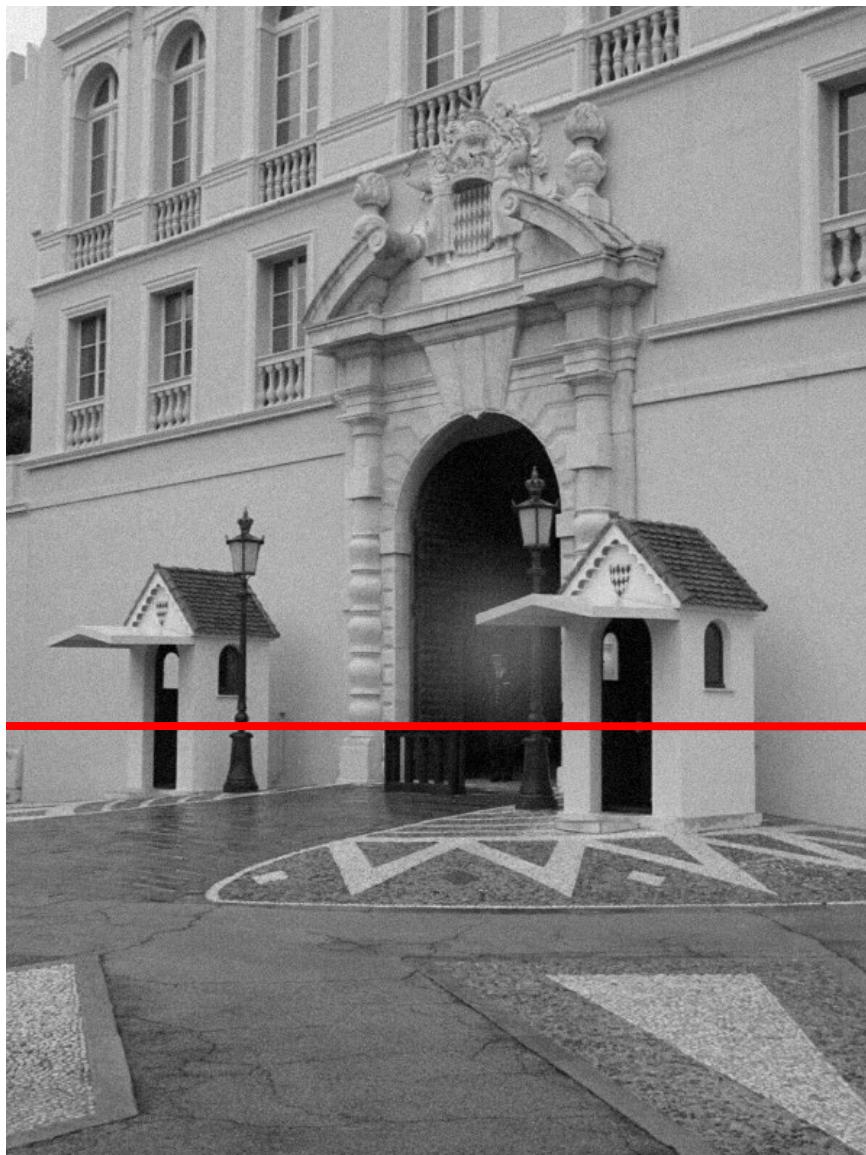


Intensity profile



Source: D. Hoiem

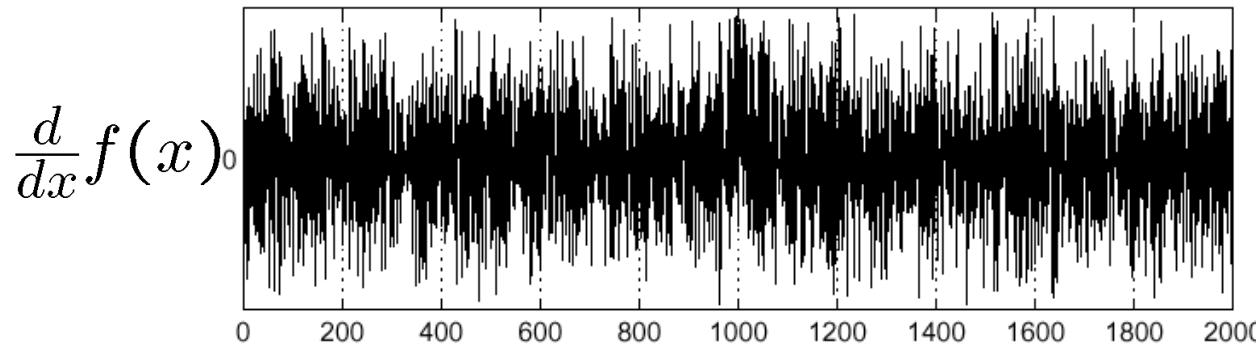
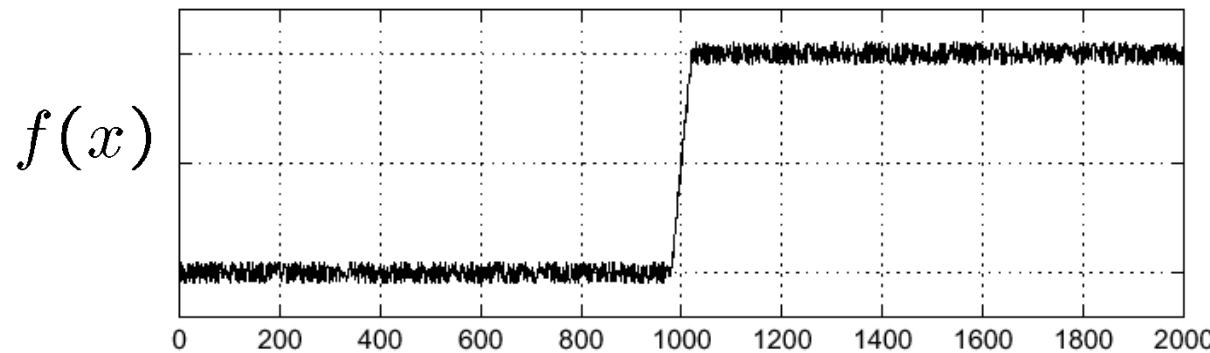
With a little Gaussian noise



Source: D. Hoiem

Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



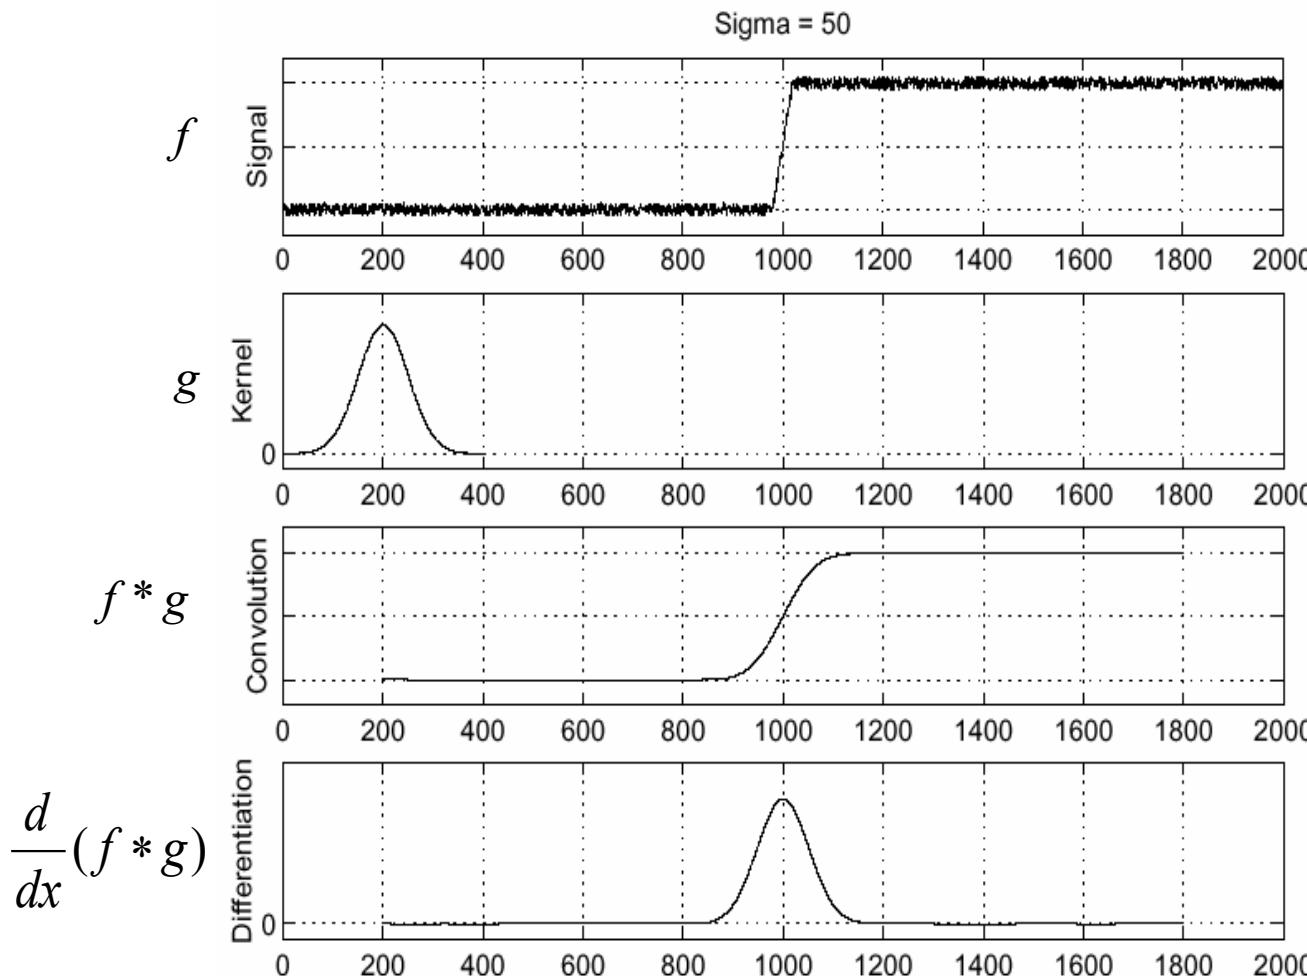
Where is the edge?

Source: S. Seitz

Effects of noise

- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What can we do about it?

Solution: smooth first

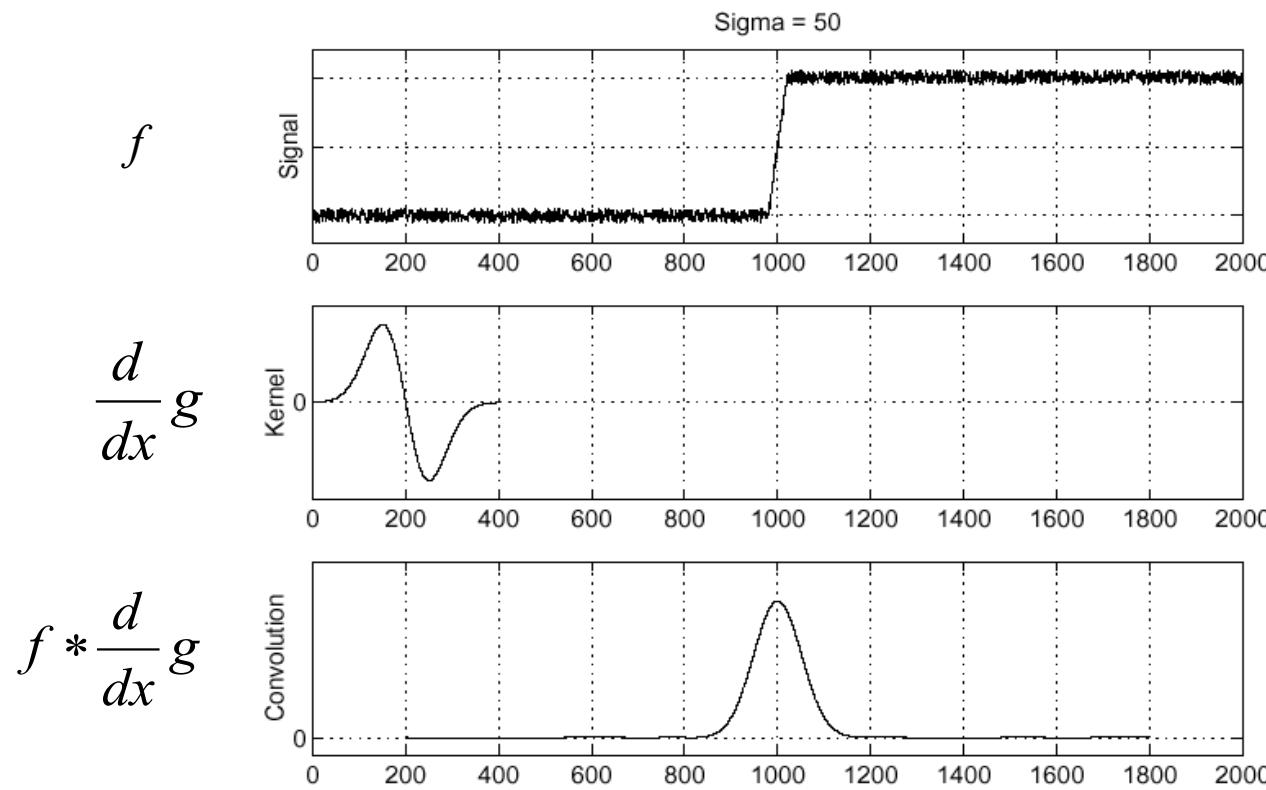


- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Source: S. Seitz

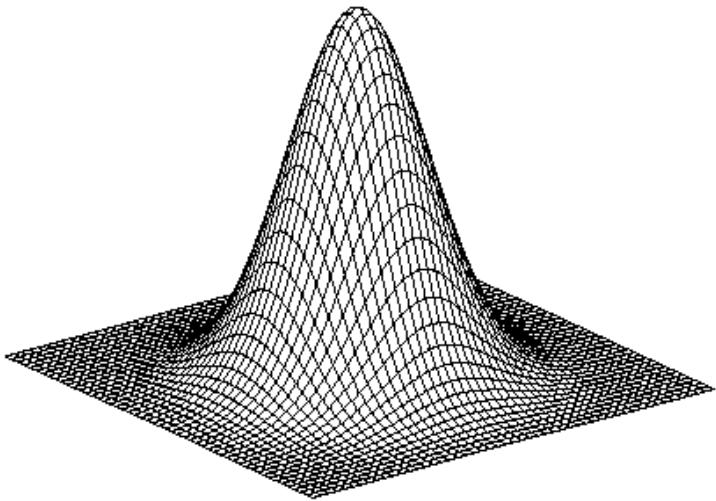
Derivative theorem of convolution

- Differentiation is convolution, and convolution is associative:
$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$
- This saves us one operation:

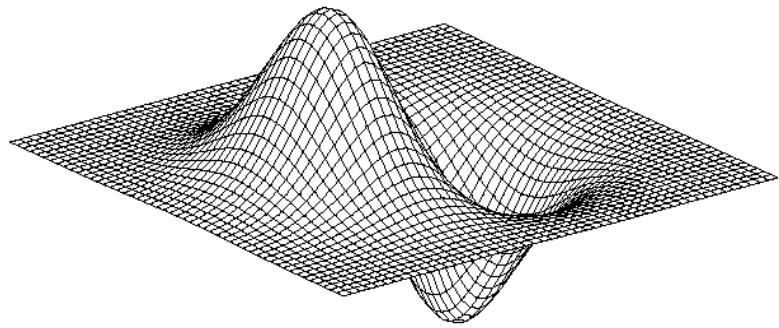


Source: S. Seitz

Derivative of Gaussian filter



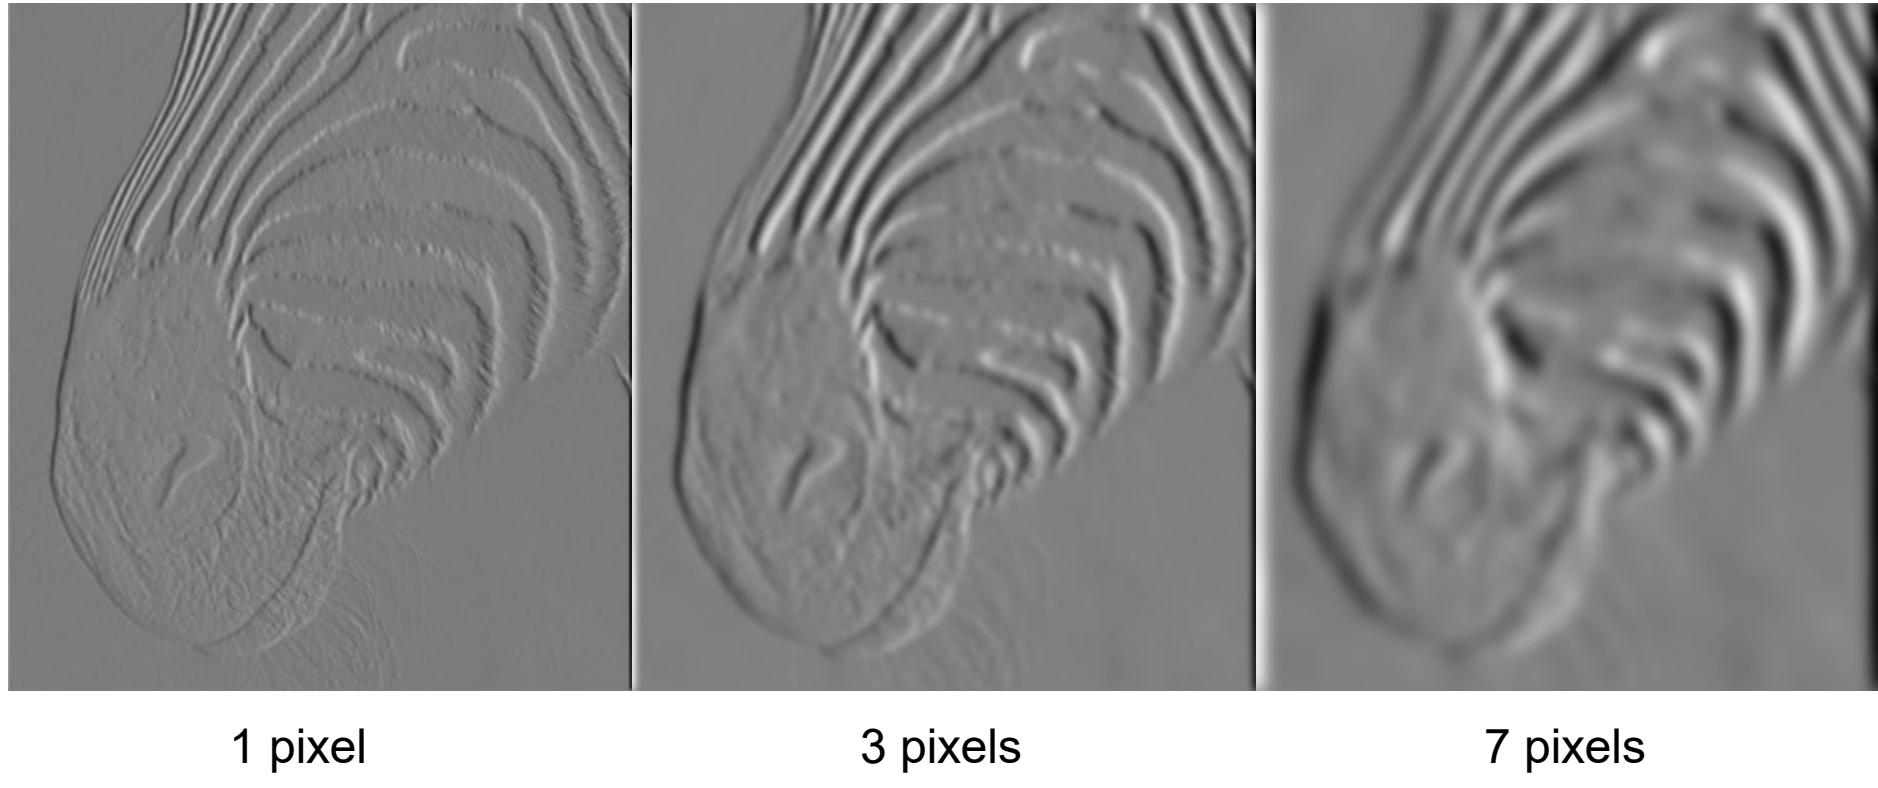
* [1 -1] =



`ldisp(kgauss(5))`

`ldisp(kdgauss(5))`

Tradeoff between smoothing and localization



1 pixel

3 pixels

7 pixels

- Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”.

Implementation issues

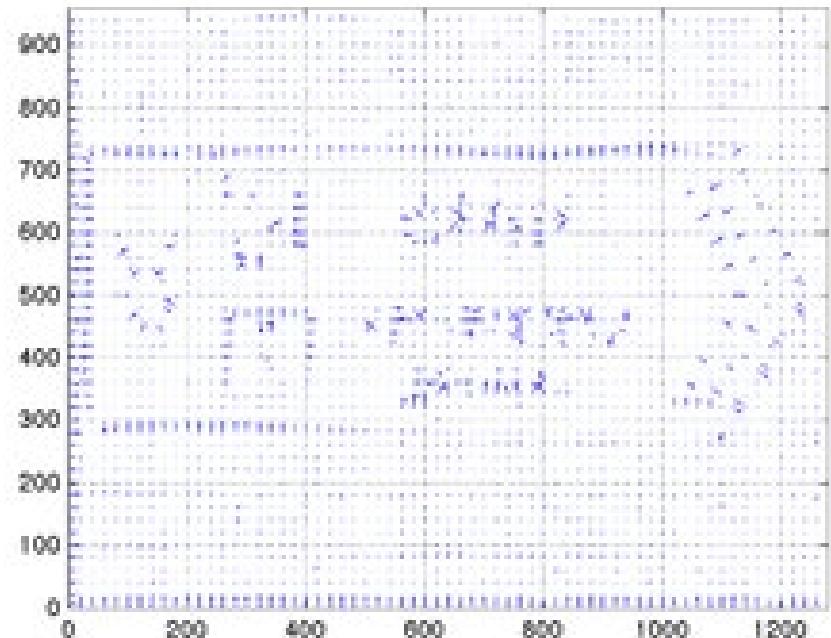
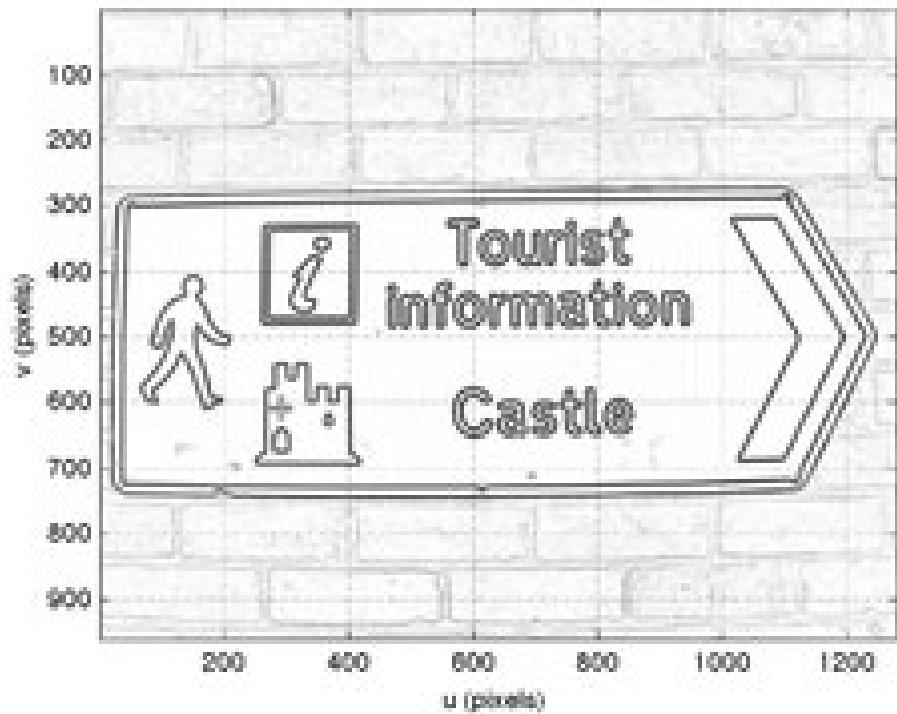


- The gradient magnitude is large along a thick “trail” or “ridge,” so how do we identify the actual edge points?
- How do we link the edge points to form curves?

Matlab vision toolbox functions- Corke

- Linear convolution
 - icanny
 - isobel
- Also can be viewed as kernels
 - Klaplace
 - Ksobel
 - kcircle

Gradient magnitude and direction



Designing an edge detector

- Criteria for a good edge detector:
 - **Good detection:** the optimal detector should find all real edges, ignoring noise or other artifacts
 - **Good localization**
 - the edges detected must be as close as possible to the true edges
 - the detector must return one point only for each true edge point
- Cues of edge detection
 - Differences in color, intensity, or texture across the boundary
 - Continuity and closure
 - High-level knowledge